

Stochastic Optimisation - Notes

Dom Hutchinson

November 23, 2020

NOTE - *Markov Chain* typically refers to the discrete setting; whilst *Markov Process* typically refers to the continuous setting.

Contents

1	Multi-Armed Bandit	3
1.1	The Problem	3
1.2	Regret Minimisation	4
1.3	Best Arm Identification for Bernoulli Distribution	5
1.4	Heuristic	5
1.5	UCB Algorithm	7
1.5.1	Analysis	8
1.5.2	Improvements?	10
1.6	Thompson Sampling	12
1.6.1	Analysis	14
1.7	Genie Analysis	16
2	Stochastic Dynamic Optimisation	17
2.1	Introduction	17
2.1.1	Induced Stochastic Process	20
2.2	Markov Decision Processes	22
2.2.1	Problem Formulation	23
2.2.2	Finite Horizon Problems	26
3	Probability	33
3.1	Probability Inequalities	34
3.2	Markov Processes	36
3.2.1	Discrete Time Markov Chains	37
3.2.2	Continuous Time Markov Process	39
3.2.3	Poisson Process	41
3.3	Transformation of Random Variables	42
0	Reference	44

0.1	Notation	45
0.1.1	Asymptotic Notation	45

1 Multi-Armed Bandit

1.1 The Problem

Example 1.1 - Motivating Example

Consider having a group of patients and several treatments they could be assigned to. How best do you go about determining which treatment is best? The obvious approach is to assign some of the patients randomly and then assign the rest to the best treatment, but how much evidence is sufficient? And how likely are you to choose a sub-optimal treatment?

Definition 1.1 - Multi-Armed Bandit Problem

An agent is faced with a choice of K actions. Each (discrete) time step the agent plays action i they receive a reward from the random real-valued distribution ν_i . Each reward is independent of the past. The distributions ν_1, \dots, ν_K are unknown to the agent.

In the *Multi-Armed Bandit Problem* the agent seeks to maximise a measure of long-run reward.

Remark 1.1 - Informal Definition of Multi-Armed Bandit Problem

Given a finite set of actions and a random reward for each action, how best do we learn the reward distribution and maximise reward in the long-run.

Definition 1.2 - Formal Definition of Multi-Armed Bandit Problem

Consider a sequence of (unknown) mutually independent random variables $\{X_i(t)\}_{i \in [1, K]}$, with $t \in \mathbb{N}$. Consider $X_i(t)$ to be the distribution of rewards an agent would receive if they performed action i at time t . Since the rewards are independent of the past $X_i(t), X_i(t+1), \dots$ are IID random variables. The *Multi-Armed Bandit Problem* tasks us to find the greatest expected reward from all the actions.

$$\mu^* := \max_{i=1}^K \mu_i \quad \text{where } \mu_i = \mathbb{E}(X_i(t))$$

There are a number of ways to formalise this objective.

Definition 1.3 - Strategy, $I(\cdot)$

Our agent's strategy $I : \mathbb{N} \rightarrow [1, K]$ is a function which determines which action the agent shall make at a given point in time. The strategy can use the knowledge gained from previous actions & their rewards only.

$$I(t) = I\left(t, \underbrace{\{I(s)\}_{s \in [1, t)}}_{\text{Prev. Actions}}, \underbrace{\{X_{I(s)}(s)\}_{s \in [1, t)}}_{\text{Prev. Rewards}}\right) \in [1, K]$$

Definition 1.4 - Long-Run Average Reward Criterion, X_*

For a strategy $I(\cdot)$ we define the following measure for *Long-Run Average Reward*

$$X_* = \liminf_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}(X_{I(t)})$$

The *Infimum* is taken as there is no guarantee the limit exists (depending on the strategy), typically we will only deal with strategies where this limit exists.

Most strategies as based only on realisations of $\{X_i(s)\}_{s \in [1, t]}$, thus $\mathbb{E}(X_{I(t)}) \leq \mu^*$ and thus $X_* \leq \mu^*$. A strategy $I(\cdot)$ is *Optimal* if $X_* = \mu^*$.

Remark 1.2 - *It is not hard to find an Optimal Strategy in the (very) long run, so we are going to look at Regret Minimisation First.*

Proposition 1.1 - *Mathematical Model & Assumptions for Multi-Armed Bandit Problem Model:*

- Bandit has K bernoulli arms.
- $X_i(t) \in \mathbb{R}$ is the reward obtained by played arm $i \in [1, K]$ at time step $t \in \mathbb{N}$.

Assumptions:

- $X_1(\cdot), X_2(\cdot), \dots$ are mutually independent sequences.
- For each i $\{X_i(t)\}_{t \in \mathbb{N}}$ is a sequence of iid $\text{Bern}(\mu_i)$ random variables

We define the following quantities to make analysis easier

- $I(t) \in [1, K]$. The index of the arm played in time t ;
- $N_j(t) := \sum_{s=1}^t \mathbb{1}(I(s) = j)$. The number of times arm j has been played in the first t rounds;
- $S_j(t) := \sum_{s=1}^t X_j(s) \mathbb{1}(I(s) = j)$. The total reward from arm j in the first t rounds. This is a Binomial random variable independent $\text{Bin}(N_j(t), \mu_j)$;
- $\hat{\mu}_{j,n} := \frac{S_j(t)}{N_j(t)}$. The sample mean reward from arm j in the first n plays of arm j .

Definition 1.5 - *Policy*

A *policy* is a family of functions f_t which specify what arm is to be played in round t . f_t should depend on the information available at time t $\{I(s), X_{I(s)}(s) : s \in [1, t-1]\}$.

Randomised policies are allowed. So, in addition to the history up to time t , f_t can depend upon a $U(t) \sim U[0, 1]$ random variable which is independent of $X_i(\cdot)$. Thus

$$I(t) = f_t \left(\underbrace{I(1), \dots, I(t-1)}_{\text{arms chosen}}, \underbrace{X_{I(1)}(1), \dots, X_{I(t-1)}(t-1)}_{\text{observed rewards}}, \underbrace{U(t)}_{\text{randomness}} \right)$$

We want to find the best policy (ie one which minimises the regret)

1.2 Regret Minimisation

Definition 1.6 - *Regret, R_n*

Regret is a measure of how much reward was lost during the first n time steps. The *Regret* R_n of a strategy $\{I(t)\}_{t \in \mathbb{N}}$ in the first n time steps is given by

$$\begin{aligned} R_n &= \max_{k=1}^K \sum_{t=1}^n \mathbb{E} \left[\underbrace{X_k(t)}_{\text{Best Pos}} - \underbrace{X_{I(t)}(t)}_{\text{Actual}} \right] \\ &= n\mu^* - \sum_{t=1}^n \mathbb{E} [X_{I(t)}(t)] \end{aligned}$$

Regret only involves expectation and thus can be learnt from observations. We want to produce a strategy where *Total Regret* grows sub-linearly. (i.e. $R_T/T \xrightarrow{T \rightarrow \infty} 0$)

Remark 1.3 - Minimising the growth rate of R_T with T is quite hard.

The best achievable regret scales as $R_T \sim c \log T$ (i.e. $R_T / c \log T \xrightarrow{T \rightarrow \infty} 1$) where c depends on the reward distributions $X_1(t), \dots, X_K(t)$.

Definition 1.7 - *Pseudo-Regret, \tilde{R}_n*

Pseudo-Regret \tilde{R}_n is a less popular alternative to *Regret R_n* . The *Pseudo-Regret \tilde{R}_n* of a strategy $\{I(t)\}_{t \in \mathbb{N}}$ in the first n time steps is given by

$$\tilde{R}_n = \max_{k=1}^K \sum_{t=1}^n (X_k(t) - X_{I(t)}(t))$$

Pseudo-Regret includes intrinsic randomness (which is independent of the past) and thus cannot be learnt from observations.

1.3 Best Arm Identification for Bernoulli Distribution

Example 1.2 - *Best Arm Identification for Bernoulli Bandits*

Consider a bandit with two *Bernoulli* arms: $\{X_1(t)\}_{t \in \mathbb{N}}$ IID RVs with distribution $\text{Bern}(\mu_1)$; and, $\{X_2(t)\}_{t \in \mathbb{N}}$ IID RVs with distribution $\text{Bern}(\mu_2)$.

Suppose $\mu_1 > \mu_2$ (i.e. arm 1 is better). Let the player play each arm n times and declare the arm with the greatest empirical mean to be the better arm. *What is the probability of choosing the wrong arm (Arm 2)?*

An error occurs if $\sum_{t=1}^n X_2(t) \geq \sum_{t=1}^n X_1(t)$ and thus we want to calculate the probability of this event.

Define $\{Y(t)\}_{t \in \mathbb{N}}$ st $Y(t) := X_2(t) - X_1(t)$. This means $Y(t) \in \{-1, 0, 1\} \subset [-1, 1]$.

To use *Hoeffding's inequality* we need to scale Y to be in $[0, 1]$, so we define $Z(t) := \frac{1}{2}(Y(t) + 1)$. We have $\mathbb{E}(Z(t)) = \frac{1}{2}(1 + \mu_2 - \mu_1)$ and an error occurs if $\sum_{t=1}^n Y(t) > 0 \iff \sum_{t=1}^n Z(t) \geq \frac{n}{2}$. By *Hoeffding's Inequality*

$$\begin{aligned} \mathbb{P}(\text{error}) &= \mathbb{P}\left(\sum_{i=1}^n Z(t) \geq \frac{n}{2}\right) \\ &= \mathbb{P}\left(\left(\sum_{i=1}^n Z(t)\right) - \frac{n}{2}(1 + \mu_2 - \mu_1) \geq \frac{n}{2}(\mu_1 - \mu_2)\right) \quad \text{subtracting } \mu \text{ from both sides} \\ &= \mathbb{P}\left(\sum_{i=1}^n \left(X_i - \underbrace{\frac{1}{2}(1 + \mu_2 - \mu_1)}_{\mu}\right) \geq n \underbrace{\frac{1}{2}(\mu_1 - \mu_2)}_t\right) \quad \text{arranging for Hoeffding's} \\ &\leq \exp\left(-2n \cdot \frac{1}{4}(\mu_1 - \mu_2)^2\right) \quad \text{by Hoeffding's Inequality} \\ &= \exp\left(-\frac{n}{2}(\mu_1 - \mu_2)^2\right) \end{aligned}$$

1.4 Heuristic

Remark 1.4 - *How many tests?*

Suppose an agent is comparing two arms and is given a finite time horizon T after in which they must choose the best arm. The obvious strategy is to perform each task N times and then choose the arm with the greatest empirical mean. But, how do we choose N to minimise

regret over time T ?

Proposition 1.2 - Naïve Heuristic (Single Test)

Consider a 2-armed bandit & the following Heuristic

Play each arm once. Pick the arm with the greatest sample mean reward (breaking ties arbitrarily) and playing that arm on all subsequent rounds.

This heuristic picks the wrong arm with probability $\mu_2(1 - \mu_1)$. In this case the wrong arm is played $T - 1$ times, giving a bounded regret

$$\mathcal{R}(T) \geq \underbrace{\mu_2(1 - \mu_1)}_{\text{prob of wrong choice}} \cdot \underbrace{(\mu_1 - \mu_2)}_{\text{Loss}} \cdot \underbrace{(T - 1)}_{\text{\# steps}}$$

This regret grows linearly in T .

Theorem 1.1 - Chernoff Bound of a Binomial Random Variable

Let $X \sim \text{Bin}(n, \alpha)$ with $n \in \mathbb{N}$, $\alpha \in (0, 1)$. Then

$$\forall \beta > \alpha \quad \mathbb{P}(X \geq \beta n) \leq e^{-nK(\beta; \alpha)}$$

where

$$K(\beta; \alpha) := \begin{cases} \beta \ln\left(\frac{\beta}{\alpha}\right) + (1 - \beta) \ln\left(\frac{1 - \beta}{1 - \alpha}\right) & \text{if } \beta \in [0, 1] \\ +\infty & \text{otherwise} \end{cases}$$

with $x \ln(x) := 0$ if $x = 0$.

Similarly

$$\forall \beta < \alpha \quad \mathbb{P}(X \leq \beta n) \leq e^{-nK(\beta; \alpha)}$$

Note that $K(\cdot; \cdot)$ is known as both *relative entropy* and *Kullback-Leibler Divergence*

Proposition 1.3 - Better Heuristic (N Tests)

Consider a 2-armed bandit problem & the following heuristic

Play each arm $N < \frac{T}{2}$. Pick the arm with the greatest sample mean reward (breaking ties arbitrarily) and playing that arm on all subsequent rounds.

Note that $S_1(n)$ & $S_2(n)$ are *binomial* random variables with distributions $\text{Bin}(N, \mu_1)$, $\text{Bin}(N, \mu_2)$ respectively. And, $S_1(n)$ and $S_2(n)$ are independent of each other. Thus for $\beta \in (\mu_2, \mu_1)$

$$\mathbb{P}(S_1(N) < \beta N, S_2(N) > \beta N) \leq e^{-N(K(\beta; \mu_1) + K(\beta; \mu_2))} = e^{-NJ(\mu_1, \mu_2)}$$

where

$$J(\mu_1, \mu_2) = \inf_{\beta \in [\mu_2, \mu_1]} (K(\beta; \mu_1) + K(\beta; \mu_2))$$

The values of β which solve $J(\cdot; \cdot)$ describe the most likely ways for the event $(S_1(N) < S_2(N))$ to occur (ie the wrong decision is made).

Proposition 1.4 - Optimal N

For the situation described in Proposition 1.2 we want to find N which minimises regret, given a total time horizon of T .

If the right decision is made in the end, regret only occurs during exploration and is equal to $N \cdot (\mu_1 - \mu_2)$ (since the wrong arm is played N times).

However, if the wrong decision is made in the end, regret is equal to $(T - N) \cdot (\mu_1 - \mu_2)$.

Thus, the overall regret up to time T is

$$\begin{aligned} \mathcal{R}(T) &= \underbrace{(T - 2N)(\mu_1 - \mu_2)\mathbb{P}(S_1(N) < S_2(N))}_{\text{if wrong decision made}} + \underbrace{N(\mu_1 - \mu_2)}_{\text{guaranteed regret}} \\ &\simeq (\mu_1 - \mu_2)(N + Te^{-NJ(\mu_1, \mu_2)}) \end{aligned}$$

This expression is minimised for N close to the solution of $1 = TJ(\mu_1, \mu_2)e^{-NJ(\mu_1, \mu_2)}$ (ie when $N = \frac{\ln T}{J(\mu_1, \mu_2)} + O(1)$).

The corresponding regret is

$$\mathcal{R}(T) = \frac{\mu_1 - \mu_2}{J(\mu_1, \mu_2)} \ln(T) + O(1)$$

If $\mu_1 \simeq \mu_2$ then $J(\mu_1, \mu_2) \simeq (\mu_1 - \mu_2)^2$ and the above regret becomes $\mathcal{R}(T) = \frac{\ln(T)}{\mu_1 - \mu_2} + O(1)$.

1.5 UCB Algorithm

Remark 1.5 - UCB Algorithm

The *Upper Confidence Bound Algorithm* is a *frequentist* algorithm for solving the multi-armed bandit problem.

Remark 1.6 - Motivation

The problem with the heuristics in **Proposition 1.2, 1.3** is that they treat the sample mean as the true mean (*Certainty Equivalence*), which is not great.

Suppose we observed sample mean reward for arm i of $\hat{\mu}_{i,n}$ after n plays. How far from the true value can μ_i be?

$$\mathbb{P}(\mu_i > \hat{\mu}_{i,n} + x) \leq e^{-2nx^2} \text{ by Hoeffding's Inequality}$$

Suppose the inequality holds with equality (ie greatest possible probability). Then for some chosen $\delta \in [0, 1]$

$$x = \sqrt{\frac{1}{2n} \ln \left(\frac{1}{\delta} \right)} \implies \mathbb{P}(\mu_i > \hat{\mu}_{i,n} + x) = \delta \quad \text{since } \delta = e^{-2nx^2}$$

This suggests a heuristic:

$$\text{Play arm which maximises } \hat{\mu}_{i, N_i(t)} + \sqrt{\frac{1}{2N_i(t)} \ln \left(\frac{1}{\delta} \right)}$$

where you choose $\delta \in [0, 1]$ based on how lucky you feel. This quantity is the upper bound of a $1 - \delta$ confidence interval for the value of μ_i .

This heuristic allows for our choice to be changed any number of times.

Definition 1.8 - UCB(α) Algorithm

Consider the set up of the multi-armed bandit problem in **Proposition 1.1** and wlog that $\mu_1 > \mu_2 \geq \dots \geq \mu_K$.

Consider a k -armed bandit and let $\alpha > 0$.

- i). In the first K rounds, play each arm once.
- ii). At the end of each round $t \geq K$ compute the $UCB(\alpha)$ index of each arm i defined as

$$\hat{\mu}_{i, N_i(t)} + \sqrt{\frac{\alpha \ln(t)}{2N_i(t)}}$$

- iii). In round $t + 1$ play the arm with the greatest index (breaking ties arbitrarily)

$$I(t + 1) = \operatorname{argmax}_{i \in [1, K]} \left\{ \hat{\mu}_{i, N_i(t)} + \sqrt{\frac{\alpha \ln(t)}{2N_i(t)}} \right\}$$

1.5.1 Analysis

Theorem 1.2 - Upper Bound on Regret

Consider a K -armed bandit and define $\Delta_i := \mu_1 - \mu_i$.

If the $UCB(\alpha)$ algorithm is used, with $\alpha > 1$, then the regret in the first T rounds is bounded above by

$$\mathcal{R} \leq \sum_{i=2}^K \left(\frac{\alpha + 1}{\alpha - 1} \Delta_i + \frac{2\alpha}{\Delta_i} \ln(T) \right)$$

This bounds grows logarithmically in T , which is very good.

If α is taken to be large, then the regret grows faster (bad). If α is small, the constant term dominates for smaller values of T (constant term blows up close to 1).

You should choose a value a bit larger than 1 (often $\alpha = 2$).

NOTE this is proved at the end of this subsection **Proof 1.4**.

Theorem 1.3 - When a sub-optimal arm is played

Consider apply $UCB(\alpha)$ to a k -armed bandit and define $\Delta_i := \mu_1 - \mu_i$. Let $s \geq K$ (so we have completed the first stage of UCB) and suppose $I(s + 1) = j \neq 1$ (ie arm at time $s + 1$ is suboptimal). Then one of the following is true:

- i). $\hat{\mu}_{1, N_1(s)} \leq \mu_1 - \sqrt{\frac{\alpha \ln(s)}{2N_1(s)}}$. The sample mean reward on the optimal arm is much smaller than the true mean.
- ii). $\hat{\mu}_{j, N_j(s)} \geq \mu_j + \sqrt{\frac{\alpha \ln(s)}{2N_j(s)}}$. The sample mean reward on arm j is much larger than its true mean.
- iii). $N_j(s) < \frac{2\alpha \ln(s)}{\Delta_j^2}$. Arm j has been played very few times.

Proof 1.1 - Theorem 1.3

This is a proof by contradiction.

Suppose $I(s+1) = j \neq 1$ but that none of the three inequalities holds. Then

$$\begin{aligned}
\underbrace{\hat{\mu}_{1,N_1(s)} + \sqrt{\frac{\alpha \ln(s)}{2N_1(s)}}}_{\text{UCB}(\alpha) \text{ index 1}} &> \mu_1 && \text{by not i)} \\
&= \mu_j + \Delta_j && \text{by def. of } \Delta_j \\
&\geq \mu_j + \sqrt{\frac{2\alpha \ln(s)}{N_j(s)}} && \text{by not iii)} \\
&\geq \hat{\mu}_{1,N_1(s)} - \sqrt{\frac{\alpha \ln(s)}{2N_1(s)}} + \sqrt{\frac{2\alpha \ln(s)}{N_j(s)}} && \text{by not ii)} \\
&\geq \hat{\mu}_{1,N_1(s)} + \left(\sqrt{2} - \frac{1}{\sqrt{2}}\right) \sqrt{\frac{\alpha \ln(s)}{N_1(s)}} \\
&= \underbrace{\hat{\mu}_{j,N_j(s)} + \sqrt{\frac{\alpha \ln(s)}{2N_j(s)}}}_{\text{UCB}(\alpha) \text{ index j}}
\end{aligned}$$

But, this implies that the $\text{UCB}(\alpha)$ index of arm 1 at the end of round s is greater than that of arm j . Hence arm j would not be played in time slot $s+1$. \square

Theorem 1.4 - Counting Lemma

Let $\{I(t)\}_{t \in \mathbb{N}}$ be a $\{0, 1\}$ -valued sequence and $N(t) := \sum_{s=1}^t I(s)$. Then

$$\forall t, u \in \mathbb{N} \quad N(t) \leq u + \sum_{s=u+1}^t I(s) \mathbb{1}\{N(s-1) \geq u\}$$

with an empty sum defined to be zero.

Proof 1.2 - Theorem 1.4

Fix $t, u \in \mathbb{N}$. There are two possibilities

Case 1 $N(t) \leq u$. (Have not reached u yet)

Case 2 $\exists s \in [1, t]$ st $N(s) > u$. (Already reached u). Let s^* denote the smallest such s . Then it must be true that $N(s^* - 1) = u$ and $s^* \geq u + 1$. Hence

$$\begin{aligned}
N(t) &= \sum_{s=1}^{s^*-1} I(s) + \sum_{s=s^*}^t I(s) \\
&= N(s^* - 1) + \sum_{s=s^*}^t I(s) \underbrace{\mathbb{1}\{N(s-1) \geq u\}}_{\text{true for all in sum}} \\
&\leq u + \sum_{s=u+1}^t I(s) \mathbb{1}\{N(s-1) \geq u\} \quad \text{since } s^* \geq u+1
\end{aligned}$$

\square

Proof 1.3 - Theorem 1.2

Fix $t \in \mathbb{N}$ and take $u_{t,j} := \left\lceil \frac{2\alpha \ln(t)}{\Delta_j^2} \right\rceil$.

By Theorem 1.4 we have that

$$N_j(t) \leq u + \sum_{s=u+1}^t \mathbb{1}\{N_j(s-1) \geq u_{t,j} \text{ \& } (I(s) = j)\}$$

Both sides involve random variables. Taking expectations we get

$$\mathbb{E}[N_j(t)] \leq u + \sum_{s=u}^{t-1} \mathbb{P}\{(N_j(s) \geq u_{y,j}) \ \& \ (I(s+1) = j)\}$$

By **Theorem 1.3** and the definition of u , *IF* $I(s+1) = j$ and $N_j(s) \geq u$ then

$$\hat{\mu}_{1,N_1(s)} \leq u_1 - \sqrt{\frac{\alpha \ln(s)}{2N_1(s)}} \quad \text{or} \quad \hat{\mu}_{j,N_j(s)} > \mu_j + \sqrt{\frac{\alpha \ln(s)}{2N_j(s)}}$$

Thus

$$\mathbb{E}[N_j(t)] \leq u_{t,j} + \sum_{s=u_{t,j}}^{t-1} \left[\underbrace{\mathbb{P}\left(\hat{\mu}_{1,N_1(s)} \leq \mu_1 - \sqrt{\frac{\alpha \ln(s)}{2N_1(s)}}\right)}_{\hat{\mu}_1 \text{ is unusually small}} + \underbrace{\mathbb{P}\left(\hat{\mu}_{j,N_j(s)} > \mu_j + \sqrt{\frac{\alpha \ln(s)}{2N_j(s)}}\right)}_{\hat{\mu}_j \text{ is unusually large}} \right]$$

By *Hoeffding's Inequality*

$$\begin{aligned} \mathbb{E}[N_j(t)] &\leq u + \sum_{s=u}^{t-1} 2s^{-\alpha} \\ &\leq u + \int_{u-1}^{\infty} 2s^{-\alpha} ds \quad \text{assumption } \alpha > 1 \text{ required here} \\ &= u + \frac{2(u-1)^{-(\alpha-1)}}{\alpha-1} \\ &\leq u + \frac{2}{\alpha-1} \quad \text{since } u \geq 2 \implies (u-1)^{-(\alpha-1)} \leq 1 \end{aligned}$$

Thus

$$\forall j \in [2, K] \quad \mathbb{E}[N_j(t)] \leq u + \frac{2}{\alpha-1} \leq \frac{2\alpha \ln(t)}{\Delta_j^2} + 1 + \frac{2}{\alpha-1}$$

A regret of $\Delta_j := \mu_1 - \mu_j$ is incurred every time arm j is played. Hence the total regret up to time t is bounded by

$$\begin{aligned} \mathcal{R}(t) &:= \sum_{i=2}^K \Delta_i \mathbb{E}[N_i(t)] \\ &\leq \sum_{i=2}^K \left(\frac{2\alpha \ln(t)}{\Delta_i} + \frac{\alpha+1}{\alpha-1} \Delta_i \right) \end{aligned}$$

□

1.5.2 Improvements?

Remark 1.7 - *The regret of UCB grows logarithmically with T . No other algorithm can do better.*

Further, the constant factor of $\ln(T)$ used is almost optimal. This shall now be shown.

Definition 1.9 - *Strongly Consistent*

A strategy for the multi-armed bandit problem is said to be *strongly consistent* if its regret satisfies $\mathcal{R}(T) = o(T^\alpha)$ for all $\alpha > 0$. (i.e. its regret grows slower than any fractional power of T).

The $UCB(\alpha)$ algorithm is strongly consistent for all $\alpha > 1$ as its regret grows logarithmically with T .

Theorem 1.5 - Lai and Robbins

Consider a K -armed bandit, where the rewards from arm i are iid $\text{Bern}(\mu_i)$ and rewards from distinct arms are mutually independent. Then, for any *strongly consistent* strategy, the number of times that a sub-optimal arm i is played up to time T , $N_i(T)$ satisfies

$$\liminf_{T \rightarrow \infty} \frac{\mathbb{E}[N_i(T)]}{\ln(T)} \geq \frac{1}{K(\mu_i; \mu^*)}$$

where $\mu^* := \max_{i=1}^K \mu_i$ and $K(q; p)$ is the *KL-Divergence* of a $\text{Bern}(q)$ distribution wrt a $\text{Bern}(p)$ distribution.

Proposition 1.5 - Lower bound on Regret

Here we derive a lower bound for the regret of any strongly consistent strategy from the multi-armed bandit problem.

$$\begin{aligned} \liminf_{T \rightarrow \infty} \frac{\mathcal{R}(T)}{\ln(T)} &= \liminf_{T \rightarrow \infty} \frac{\sum_{i; \mu_i < \mu^*} (\mu^* - \mu_i) \mathbb{E}[N_i(T)]}{\ln(T)} \\ &\geq \sum_{i; \mu_i < \mu^*} \frac{\mu^* - \mu_i}{K(\mu_i; \mu^*)} \end{aligned} \quad \text{by Theorem 1.6}$$

Proposition 1.6 - Comparing to Lower bound of $UCB(\alpha)$

We showed that the regret of the $UCB(\alpha)$ algorithm satisfies

$$\limsup_{T \rightarrow \infty} \frac{\mathcal{R}(T)}{\ln(T)} \leq \sum_{i; \mu_i < \mu^*} \frac{2}{\mu^* - \mu_i}$$

To compare this to the result in **Proposition 1.5** we use *Pinsker's Inequality*. (Proof in homework).

We see thjat the upper bound on the regret achieved by $UCB(\alpha)$ is approximately four times greater than the lower bound on the best regret achievable by any algorithm. This is very good.

Theorem 1.6 - Concentration Inequalities for Sample Means

$$\begin{aligned} \mathbb{P} \left(\hat{\mu}_{j, N_j(s)} \geq u_j + \sqrt{\frac{\alpha \ln s}{2N_j(s)}} \right) &\leq e^{-\alpha \ln s} = s^{-\alpha} \\ \mathbb{P} \left(\hat{\mu}_{1, N_1(s)} \leq u_1 - \sqrt{\frac{\alpha \ln s}{2N_1(s)}} \right) &\leq e^{-\alpha \ln s} = s^{-\alpha} \end{aligned}$$

Proof 1.4 - Theorem 1.5

The proof is immediate from *Hoeffding's Inequality*, which is applicable since the X_j are iid and take values in $\{0, 1\} \subseteq [0, 1]$.

Remark 1.8 - Is there an algorithm which achieves lower regret?

No. There is no algorithm which has regret growing slower than $\ln(T)$.

1.6 Thompson Sampling

Remark 1.9 - Thompson Sampling

Thompson Sampling is a *Bayesian* algorithm for the multi-armed bandit problem. It was one of the first algorithms for solving the problem, but remains one of the best as it is asymptotically optimal.

Theorem 1.7 - Relationship between Beta & Gamma Distribution

Let $X \sim \text{Gamma}(\alpha, \lambda)$ and $Y \sim \text{Gamma}(\beta, \lambda)$ (ie shared scale parameter but different shape parameters). Then

$$V := \frac{X}{X+Y} \sim \text{Beta}(\alpha, \beta)$$

Proof 1.5 - Theorem 1.7

Consider the map $(X, Y) \mapsto \frac{X}{X+Y}$. This maps $\mathbb{R}^2 \rightarrow \mathbb{R}$ (reduces dimensions) and thus we cannot directly use the formula to compute the density of V .

We introduce an auxiliary random variable $W = X + Y$. Now $(X, Y) \mapsto (V, W) := g(X, Y) = \left(\frac{X}{X+Y}, X+Y \right)$. Note that X and Y are non-negative random variables. Hence, $V \in [0, 1]$ and $W \in \mathbb{R}^+$.

For $(v, w) \in [0, 1] \times \mathbb{R}^+$ we want to find all solutions of $g(x, y) = (v, w)$. Clearly $(x, y) = \left(w, w \frac{1-v}{v} \right)$ is a unique solution. The joint density of (V, W) is given by the formula

$$f_{V,W}(v, w) = \sum_{(x,y):g(x,y)=(v,w)} \frac{f_{X,Y}(x, y)}{|\det(J_g(x, y))|} \quad (1)$$

Using the density function for *Gamma* random variables and the independence of X and Y

$$f_{X,Y}(x, y) = \frac{1}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} y^{\beta-1} e^{-x-y}$$

Next, we compute the *Jacobian* of g and its determinant

$$J_g(x, y) = \begin{pmatrix} \frac{y}{(x+y)^2} & \frac{-x}{(x+y)^2} \\ 1 & 0 \end{pmatrix} \implies |\det(J_g(x, y))| = \frac{x}{(x+y)^2} = \frac{v^2}{w}$$

Substituting the definition of $f_{X,Y}$, the results of the *Jacobian* and the solution $(x, y) = \left(w, w \frac{1-v}{v} \right)$ into (1) we get

$$\begin{aligned} f_{V,W}(v, w) &= \frac{1}{v^2/w} \frac{1}{\Gamma(\alpha)\Gamma(\beta)} w^{\alpha-1} e^{-w} \left(\frac{w(1-v)}{v} \right)^{\beta-1} e^{-w(1-v)/v} \\ &= \frac{1}{\Gamma(\alpha)\Gamma(\beta)} \frac{w}{v^2} w^{\alpha-1} \left(w \frac{1-v}{v} \right)^{\beta-1} e^{-w/v} \end{aligned}$$

We are only interested in the marginal distribution of V

$$\begin{aligned}
f_V(v) &= \int_{w=0}^{\infty} f_{V,W}(v, w) dw \\
&= \frac{(1-v)^{\beta-1} v^{\alpha} - 1}{\Gamma(\alpha)\Gamma(\beta)} \int_{w=0}^{\infty} \left(\frac{w}{v}\right)^{\alpha+\beta-1} e^{-w/v} \frac{dw}{v} \\
&= \frac{(1-v)^{\beta-1} v^{\alpha-1}}{\Gamma(\alpha)\Gamma(\beta)} \int_{u=0}^{\infty} u^{\alpha+\beta-1} e^{-u} du \quad \text{where } u := \frac{w}{v} \\
&= \frac{\Gamma(\alpha+\beta)}{\Gamma(\alpha)\Gamma(\beta)} v^{\alpha-1} (1-v)^{\beta-1} \\
&\sim \text{Beta}(\alpha, \beta)
\end{aligned}$$

□

Remark 1.10 - *Beta(1,1) is equivalent to Uniform[0,1]*

Proposition 1.7 - *Thompson Sampling Algorithm*

Consider a K -armed bandit with independent Bernoulli arms with parameters μ_1, \dots, μ_K . *Thompson Sampling* follows the following process.

- i). Define a prior distribution $\text{Beta}(1,1)$ for the parameter of each arm.
- ii). At the start of round t , sample $\hat{\mu}_i(t)$ for $i \in [1, K]$, from the corresponding prior distributions.
- iii). Play the arm with the greatest sample value $I(t) \in \arg\max_{i \in [1, K]} \hat{\mu}_i(t)$ (breaking ties arbitrarily).
- iv). Compute a posterior distribution for that parameter, based on the observed reward. $\text{Beta}(\alpha + 1, \beta)$ if a reward is given and $\text{Beta}(\alpha, \beta + 1)$ if a reward is not given. (The priors for the other arms are the same as their priors as no result was observed).
- v). Keep repeating ii)-iv) using the posterior distributions calculated in round t as the priors for round $t + 1$.

The endpoint of the algorithm is when time runs out.

Proposition 1.8 - *Choosing the Prior Distribution for Thompson Sampling*

We use *Beta* distributions as the prior for the unknown parameters μ_i of the *Bernoulli* reward distributions. This choice is convenient because, if the prior has a *Beta* distribution then so does the posterior distribution, after observing the reward (i.e. it is a *Conjugate Prior*).

Remark 1.11 - *Beta Distributions are conjugate priors for Bernoulli Random Variables*

Remark 1.12 - *Thompson Sampling on other Reward Distributions*

Thompson Sampling can be extended to non-bernoulli reward distributions. However, the ease of implementation depends on how easy it is to sample from the posterior distribution. *Conjugate priors* should be used to make this easier. Bounds on the regret are not known in all cases.

Theorem 1.8 - *Posterior from a Beta Distribution Prior of a Bernoulli Random Variable*

Let $X \sim \text{Bernoulli}(\mu)$ and assume a prior distribution $\mu \sim \text{Beta}(\alpha, \beta)$. Then, the posterior distribution of μ given $X = 1$ is $\text{Beta}(\alpha + 1, \beta)$ and given $X = 0$ is $\text{Beta}(\alpha, \beta + 1)$.

Proof 1.6 - *Theorem 1.8*

Let f_μ denote the density of μ . By our assumption of a $\text{Beta}(\alpha, \beta)$ prior we have

$$f_\mu(p) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} p^{\alpha-1} (1-p)^{\beta-1} \quad \text{for } p \in [0, 1]$$

Also $\mathbb{P}(X = 1 | \mu = p) = p = 1 - \mathbb{P}(X = 0 | \mu = p)$.

Hence, by *Bayes's Theorem*, the posterior density conditional on $X = 1$ satisfies

$$\begin{aligned} f_\mu(p | X = 1) &\propto f_\mu(p) \mathbb{P}(X = 1 | p) \\ &\propto p^\alpha (1-p)^{\beta-1} \end{aligned}$$

Where the constant of proportionality is determined by the fact the density integrates to 1.

We can recognise this distribution as the pdf of a $\text{Beta}(\alpha + 1, \beta)$. Similarly, it can be shown that $f_\mu(\cdot | X = 0)$ is the density of a $\text{Beta}(\alpha, \beta + 1)$ random variable. \square

Theorem 1.9 -

Let $X \sim \text{Beta}(\alpha, \beta)$ for $\alpha, \beta \in \mathbb{N}$ and $Y \sim \text{Bin}(\alpha + \beta - 1, p)$ for $p \in (0, 1)$. Then

$$\mathbb{P}(X > p) = \mathbb{P}(Y \leq \alpha - 1)$$

Proof 1.7 - Theorem 1.9

Consider this result. Let $\{N_t\}_{t \geq 0}$ be a poisson process with intensity $\lambda > 0$ and $n, t \in \mathbb{N}$. Then, conditional on the event $N_t = n$, the unordered increment times on $[0, t]$ are mutually independent and uniformly distributed on $[0, t]$.

Let $X \sim \text{Beta}(\alpha, \beta)$ then we can write $X = \frac{V}{V+W}$ where $V \sim \text{Gamma}(\alpha, 1)$, $W \sim \text{Gamma}(\beta, 1)$. If α, β are integer values then we can interpret V as the time of the α^{th} increment of a unit rate Poisson process and $(V + W)$ as the time of the $(\alpha + \beta)^{\text{th}}$ increment.

Consequently, conditional on $(V+W) = \tau$, the poisson process N_t has exactly $\alpha + \beta - 1$ increments in the interval $[0, \tau)$. By the fact result above, the unordered times of these increments are iid uniformly distributed on $[0, \tau)$.

The event $\{X > p\}$ is the same as the event $\{V > p\tau\}$, conditional on $(V + W) = \tau$. This means that at most $\alpha - 1$ increments occur in $[0, p\tau]$. As the increments are IID uniform in $[0, \tau)$ and there are $\alpha + \beta - 1$ in total, the number of increments in $[0, p\tau]$ has a $\text{Bin}(\alpha + \beta - 1, p)$ distribution.

Thus the events $\{X > p\}$, $\{V > p\tau | V + W = \tau\}$ and $\{Y \leq \alpha - 1\}$ all have the same probability. \square

1.6.1 Analysis

Remark 1.13 - Analysis of Thompson Sampling is Hard

The main challenge is to deal with the situation where there is an initial run of bad luck on the optimal arm. This causes the posterior for the optimal arm to be biased towards small values. Hence, the optimal arm is not played very often meaning it takes a long time to recover from the initial bad luck.

For contrast, we only need to worry about plays of sub-optimal arms after they have been played sufficiently often, by which time the posterior is concentrated around the true parameter value.

Theorem 1.10 - Bound on Regret

The regret of *Thompson Sampling* applied to a multi-armed bandit with two *Bernoulli* arms is bounded as

$$\mathcal{R}(T) \leq \frac{40 \ln(T)}{\Delta} + c$$

where Δ is the arm gap and c is some arbitrary constant which depends on Δ (but, importantly, not T).

The proof to this is not given in full, but some useful lemmas are shown.

Remark 1.14 - Posterior Distribution over Time

As the number of times each arm is played increases, its posterior distribution concentrates increasingly sharply around the true parameter value.

Proposition 1.9 - Number of times wrong arm is played

In the analysis of a multi-armed bandit with two *Bernoulli* arms we assume that the second arm (with parameter μ_2) is the worse of the two arms. Thus to bound regret it suffices to bound the number of times the second arm is played.

Fix a time horizon T and define $L = \left\lceil \frac{24 \ln(T)}{\Delta^2} \right\rceil$, $\tau = \inf\{t \in [0, T] : N_2(t) \geq L\}$. Then, for $\tau \leq t \leq T$

$$\mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}\right) \leq \frac{2}{T^3}$$

Where $\theta_i(t)$ is the value sampled from the prior of μ_i at time t .

Proof 1.8 - Proposition 1.9

By the definition of τ if $t \geq \tau$ then $N_2(t) \geq L$. Thus

$$\begin{aligned} & \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}\right) \\ = & \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}, \frac{S_2(t)}{N_2(t)} \leq \mu_2 + \frac{\Delta}{4}\right) + \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}, \frac{S_2(t)}{N_2(t)} > \mu_2 + \frac{\Delta}{4}\right) \\ \leq & \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2} \mid \frac{S_2(t)}{N_2(t)} \leq \mu_2 + \frac{\Delta}{4}\right) + \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}, \frac{S_2(t)}{N_2(t)} > \mu_2 + \frac{\Delta}{4}\right) \quad (1) \end{aligned}$$

We now bound these two terms.

Firstly, conditional on the number of times the second arm is played $N_2(T)$, the total reward from these plays $S_2(t)$ is the sum of $N_2(t)$ independent $\text{Bern}(\mu_2)$ random variables. Hence, using *Hoeffding's Inequality* we have

$$\mathbb{P}\left(\frac{S_2(t)}{N_2(t)} > \mu_2 + \frac{\Delta}{4} \mid N_2(t)\right) \leq \exp\left(-2N_2(t) \frac{\Delta^2}{16}\right)$$

As we have assumed that $N_2(t) \geq L \geq \frac{1}{\Delta^2}(24 \ln(T))$ we can conclude that

$$\mathbb{P}\left(\frac{S_2(t)}{N_2(t)} > \mu_2 + \frac{\Delta}{4}\right) \leq \exp(-3 \ln(T)) = \frac{1}{T^3} \quad (2)$$

Next, we note that conditional on $S_2(t)$ and $N_2(t)$, the distribution of $\theta_2(t)$ is $\text{Beta}(S_2(t) + 1, N_2(t) - S_2(t) + 1)$. Consequently, by **Theorem 1.9**, we have that

$$\mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}\right) = \mathbb{P}\left(\text{Bin}\left(N_2(t) + 1, \mu_2 + \frac{\Delta}{2}\right) \leq S_2(t)\right)$$

Applying *Hoeffding's inequality* to the RH term, we see that for $N_2(t) \geq L$ we have

$$\begin{aligned} \mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2} \mid \frac{S_2(t)}{N_2(t)} \leq \mu_2 + \frac{\Delta}{4}\right) &\leq \exp\left(-2(N_2(t) + 1)\frac{\Delta^2}{16}\right) \\ &\leq \exp\left(-\frac{L\Delta^2}{8}\right) \\ &\leq \frac{1}{T^3} \end{aligned} \quad (3)$$

Substituting (2) and (3) into (1), we can conclude that if $t \geq \tau$ (ie $N_2(t) \geq L$) then

$$\mathbb{P}\left(\theta_2(t) \geq \mu_2 + \frac{\Delta}{2}\right) \leq \frac{2}{T^3}$$

□

1.7 Genie Analysis

Remark 1.15 - Genie

Here we analyse a simpler version of Thompson Sampling for a 2-armed bandit. We assume that the value of μ_1 is known, but it is not known whether it is greater than μ_2 .

This means we now only have a prior/posterior for μ_2 and we compare $\theta_2(t)$ (the value sampled from the prior of μ_2) to the true value of μ_1 .

It is likely that this scenario should be more successful than the standard scenario, thus we can only find an upper bound on the regret of the normal scenario.

Theorem 1.11 - Times Sub-optimal arm is played

Fix $T \in \mathbb{N}$ and define $L := \left\lceil \frac{2 \ln(T)}{\Delta^2} \right\rceil$, $\tau := \inf\{t \in [1, T] : N_2(t) \geq L\}$. τ is the first time that arm two has been played at least L times. The number of plays of arm two after τ is bounded as

$$\forall t \geq \tau \quad \mathbb{P}(\theta_2(t) \geq \mu_1) \leq \frac{2}{T}$$

This means $\mathbb{E}[\text{plays of arm two after time } \tau] = (T - \tau)\frac{2}{T} \leq 2$

Proof 1.9 - Theorem 1.11

Define the events

$$A_t := \{\theta_2(t) \geq \mu_1\} \quad B_t := \left\{ \frac{S_2(t)}{N_2(t)} \leq \mu_2 + \frac{\Delta}{2} \right\}$$

A_t is the event that the sample from the prior of μ_2 is greater than μ_1 . B_t is the event the observed rewards from arm 2 are closer to μ_2 than μ_1 .

If $t \geq \tau$, then $N_2(t) \geq L$ and Hoeffding's inequality yields

$$\begin{aligned} \mathbb{P}(B_t^c) &\leq \exp\left(-2N_t \frac{\Delta^2}{4}\right) \\ &\leq \exp\left(-\frac{\Delta^2 L}{2}\right) \quad \text{by def. } L \\ &\leq e^{-\ln(T)} \\ &= \frac{1}{T} \end{aligned}$$

We can bound $\mathbb{P}(A_t)$ as follows

$$\begin{aligned} \mathbb{P}(A_t) &= \mathbb{P}(A_t \cap B_t) + \mathbb{P}(A_t \cap B_t^c) \\ &= \mathbb{P}(A_t|B_t)\mathbb{P}(B_t) + \mathbb{P}(A_t|B_t^c)\mathbb{P}(B_t^c) \\ &\leq \mathbb{P}(A_t|B_t) + \mathbb{P}(B_t^c) \end{aligned}$$

We now want to bound $\mathbb{P}(A_t|B_t)$.

Since $\theta_2(t+1)$ is sampled from the posterior distribution of μ_2 after t rounds we have

$$\theta_2(t+1) \sim \text{Beta} \left(1 + \underbrace{S_2(t)}_{\# \text{ successes}}, 1 + \underbrace{N_2(t) - S_2(t)}_{\# \text{ failures}} \right)$$

Hence, by **Theorem 1.9**

$$\mathbb{P}(\theta_2(t+1) \geq \mu_1 | S_2(t), N_2(t)) = \mathbb{P}(\text{Bin}(N_2(t) + 1, \mu_1) \leq S_2(t))$$

By Hoeffding's inequality, if $S_2(t) < \mu_1 \cdot N_2(t)$, then

$$\mathbb{P}(\text{Bin}(N_2(t) + 1, \mu_1) \leq S_2(t)) \leq \exp \left(-2N_2(t) \left(\mu_1 - \frac{S_2(t)}{N_2(t)} \right)^2 \right)$$

Conditioning on the event B_t , we have

$$\begin{aligned} S_2(t) &\leq \left(\mu_2 + \frac{\Delta}{2} \right) N_2(t) \\ &= \left(\mu_1 - \frac{\Delta}{2} \right) N_2(t) \end{aligned}$$

Hence

$$\begin{aligned} \mathbb{P}(A_t|B_t, N_2(t)) &= \mathbb{P}(\theta_2(t+1) \geq \mu_1 | B_t, N_2(t)) \\ &\leq \exp \left(-\frac{2N_2(t)\Delta^2}{4} \right) \end{aligned}$$

Recall, by the definition of τ , $\forall t \geq \tau$, $N_2(t) \geq L$.

Hence, $\forall t \geq \tau$

$$\begin{aligned} \mathbb{P}(A_t|B_t) &\leq \exp \left(-\frac{L\Delta^2}{2} \right) \\ &\leq e^{-\ln T} \\ &= \frac{1}{T} \end{aligned}$$

We have already showed $\forall t \geq \tau$, $\mathbb{P}(B_t^c) \leq \frac{1}{T}$.

Combining these results, we get

$$\forall t \geq \tau, \quad \mathbb{P}(A_t) \leq \mathbb{P}(A_t|B_t) + \mathbb{P}(B_t^c) \leq \frac{2}{T}$$

This is claim on **Theorem 1.11**

□

Proposition 1.10 - Bound of Regret

Using **Theorem 1.11** above we can bound the regret as

$$\mathcal{R}(T) \leq \Delta \cdot (L + 2)$$

where $L + 2$ is the most time arm two is played in the first T time steps.

2 Stochastic Dynamic Optimisation

2.1 Introduction

Remark 2.1 - AKA

Stochastic Dynamic Programming or Sequential Decision Making

Definition 2.1 - Types of Optimisation Problem

Stochastic The quantity to be optimised involves random parameters

Dynamic Actions are taken in stages, over a period of time, so later actions have more information available to them.

Functional A *Decision Function* maps from system-states to the action to be taken.

Definition 2.2 - Stochastic Dynamic Optimisation

Stochastic Dynamic Optimisation is a branch of probability and applied mathematics which deals with the problem of making optimal (or nearly-optimal) decisions and actions in stochastic systems.

Definition 2.3 - Stochastic System

A *Stochastic System* is a dynamic system whose model involves uncertain parameters which can be represented by random variables.

Definition 2.4 - Types of Stochastic Dynamic Processes

- *Sequential Decision Process* - An agent is able to observe the state of a stochastic system and use this information to make a decision on what action to take. The action causes the system to evolve and the agent can then observe again before making their next decision. These are *Functional Optimisation Problems*.

Proposition 2.1 - Formulation of a Sequential Decision Problem

In a *Sequential Decision Problem* the agent in a *Sequential Decision Process* is tasked with choosing a sequence of actions st that the system performs optimally wrt a predetermined performance criterion.

The actions of the agent are information by the available relevant information contained in the current state, past states and past actions. Actions can be considered as functions of these sources of information.

Here are the main elements of a *Sequential Decision Problem*

- *Decision Epochs* - The periods in which decisions are taken and their effect realised.
- *System States* - The value of variables within the system at points in time.
- *Agent's Actions* - Actions taken by an agent which directly affect the system.
- *Immediate Rewards/Costs* - The reward/cost the agent receives/incurs by taking this action.
- *Mathematical Model of Stochastic System* - Collection of mathematical equations describing the system within which the process is occurring. Particularly, the affect of actions.

Example 2.1 - Motivating Example - Inventory Control Problem

Stochastic Dynamic Optimisation was born out of the *Inventory Control Problem* (A *Sequential Decision Problem*).

The single-item *Inventory Control Problem* considers the amount of a single item stored. A quantity of this item is ordered (by us) and sold (to customers) at discrete time periods in $[0, N)$ and the amount ordered is stochastic. We are tasked to meet this stochastic demand, while keeping costs incurred from ordering and storing to a minimum.

Let X_t denote the stock available at time t , Y_t be the stock ordered by us at time t , and Z_t be the demand from customers at time t . We assume that Z_0, \dots, Z_{N-1} are IID random variables with some known distribution and that excess demand is backlogged (and fulfilled as soon as additional inventory is available).

We can derive the following equation for stock levels at each time period

$$X_{t+1} = X_t + Y_t - Z_t \text{ for } t \in [1, N)$$

if $X_t < 0$ there is a backlog of orders.

Let $c > 0$ be the cost to order each item, so the total order cost at time t is (cY_t) . Let $g(|x|)$ represent the cost of holding x stock end of the time period; either as a storage cost (for $x > 0$) or a penalty for having a backlog ($x < 0$). Note that $g(\cdot)$ should be a non-negative function. The total cost over N time period is

$$\text{TotalCost} = \sum_{t=0}^{N-1} (cY_t + \underbrace{g(X_t + Y_t - Z_t)}_{\equiv X_{t+1}})$$

We want to minimise **TotalCost** by making appropriate choices of Y_0, \dots, Y_{N-1} (given $Y_i \geq 0 \forall i$).

We cannot directly minimise the expression of **TotalCost** as its value depends on the stochastic random variables Z_0, \dots, Z_{N-1} . Thus we try to minimise the expected value $\mathbb{E}[\text{TotalCost}]$

$$\mathbb{E}[\text{TotalCost}] := \mathbb{E} \left(\sum_{t=0}^{N-1} (cY_t + g(X_t + Y_t - Z_t)) \right)$$

We can take either a: Static (**Proposition 2.3**); or Dynamic Approach (**Proposition 2.4**).

Proposition 2.2 - Formulation of a Sequential Decision Problem

Here are the specification of the components of a sequential decision problem, as specified in **Proposition 2.1**.

- *Decision Epochs* - Beginnings of the time periods $0, \dots, N-1$
- *System States* - Stock levels X_0, \dots, X_{N-1}
- *Agent's Actions* - Ordering new stock Y_0, \dots, Y_{N-1}
- *Immediate Costs* - Storage costs/penalty for underfuling orders $cY_t + g(X_t + Y_t - Z_t)$
- *Mathematical Model of Stochastic System* - A single state equation $X_{t+1} = X_t + Y_t - Z_t$ and the distributions of demand Z_0, \dots, Z_{N-1} .

Proposition 2.3 - Static Approach to Example 2.1

In the *Static Approach* to the *Inventory Control Problem* Y_0, \dots, Y_{N-1} are treated as deterministic variables. We choose the values for Y_0, \dots, Y_{N-1} at the very start (before any transactions occur) and thus are chosen before stock levels X_0, \dots, X_{N-1} or sale quantities (Z_0, \dots, Z_{N-1}) are available. Thus this approach is inevitably sub-optimal.

Proposition 2.4 - Dynamic Approach to Example 2.2

Since there no penalty in not choosing the order quantity Y_t before time t , there is no need to take a static approach. Moreover, it is highly advantageous to delay choosing the order quantity until as much relevant information (e.g. Demand & Stock Levels) is known as possible.

In this *Dynamic Approach* we can consider Y_t to be a function of current and past stock levels X_0, \dots, X_t

$$Y_t = d_t(X_0, \dots, X_t)$$

We want to find *Decision Functions* $d_0(\cdot), \dots, d_{N-1}(\cdot)$ st that the *Expected Total Cost* is minimised.

2.1.1 Induced Stochastic Process

Definition 2.5 - Induced Stochastic Process $\{(X_t, Y_t)\}_{t \in T}$

An *Induced Stochastic Process* is the set of actions and states over a time horizon T , $\{(X_t, Y_t)\}_{t \in T}$. An *Induced Stochastic Process* is fully specified by

- i). The pmf of X_0 .
- ii). The transition probabilities $\{p_t(\cdot|\cdot)\}_{t \in T}$.
- iii). A decision policy $\pi := \{q_t(\cdot|\cdot)\}_{t \in T}$.

Theorem 2.1 - Marginal Distribution of Actions & States

The induced stochastic process satisfies the following equation for all time-horizons T , state-spaces S and action-spaces A

$$\mathbb{P}(X_{0:t} = s_{0:k}, Y_{0:t-1} = a_{0:t-1}) = \mathbb{P}(X_0 = s_0) \prod_{k=0}^{t-1} \underbrace{p_k(s_{k+1}|s_k, a_k)}_{\text{transition}} \underbrace{q_k(a_k|s_{0:k}, a_{0:k-1})}_{\text{decision}}$$

In the case of a deterministic decision rule the term $q_k(a_k|s_0, \dots, s_k, a_0, \dots, a_{k-1})$ should be replaced by $\mathbb{1}\{d_k(s_{0:k}, a_{0:k-1}) = a_k\}$

Proof 2.1 - Theorem 2.1

$$\begin{aligned} & \mathbb{P}^\pi(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \\ &= \mathbb{P}^\pi(X_0 = s_0) \prod_{k=1}^t \mathbb{P}^\pi(X_k = s_k, Y_{k-1} = a_{k-1} | X_{0:k-1} = s_{0:k-1}, Y_{0:k-2} = a_{0:k-2}) \text{ by chain rule} \\ &= \mathbb{P}(X_0 = s_0) \prod_{k=1}^t \frac{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})}{\mathbb{P}^\pi(X_{0:k-1} = s_{0:k-1}, Y_{0:k-2} = a_{0:k-2})} \text{ by def. of conditional} \\ &= \mathbb{P}(X_0 = s_0) \prod_{k=0}^{t-1} \frac{\mathbb{P}^\pi(X_{0:k+1} = s_{0:k+1}, Y_{0:k} = a_{0:k})}{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})} \\ &= \mathbb{P}(X_0 = s_0) \prod_{k=0}^{t-1} \left(\frac{\mathbb{P}^\pi(X_{0:k+1} = s_{0:k+1}, Y_{0:k} = a_{0:k})}{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})} \cdot \frac{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})}{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})} \right) \end{aligned}$$

I shall simplify both terms of the product separately. Consider the first

$$\begin{aligned} & \frac{\mathbb{P}^\pi(X_{0:k+1} = s_{0:k+1}, Y_{0:k} = a_{0:k})}{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})} \\ &= \frac{\mathbb{P}^\pi(X_{k+1} = s_{k+1}, X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})}{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})} \text{ by rearrangement} \\ &= \mathbb{P}^\pi(X_{k+1} = s_{k+1} | X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k}) \text{ by def. of conditional} \\ &= \mathbb{P}^\pi(X_{k+1} = s_{k+1} | X_k = s_k, Y_k = a_k) \text{ by Markov property} \\ &= p_K(s_{k+1} | s_k, a_k) \end{aligned}$$

This is the *Transition Probability* at epoch k . Now consider the second term

$$\begin{aligned}
& \frac{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k} = a_{0:k})}{\frac{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})}{\mathbb{P}^\pi(Y_k = a_k, X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})}} \\
&= \frac{\mathbb{P}^\pi(X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})}{\mathbb{P}^\pi(Y_k = a_k | X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1})} \text{ by rearrangement} \\
&= \mathbb{P}^\pi(Y_k = a_k | X_{0:k} = s_{0:k}, Y_{0:k-1} = a_{0:k-1}) \text{ by def. of conditional} = \\
& q_k(a_k | s_{0:k}, a_{0:k-1})
\end{aligned}$$

This is the *Decision Probability* at epoch k . By substituting these two simplifications into the original expression, we get

$$\mathbb{P}^\pi(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) = P(X_0 = s_0) \prod_{k=0}^{t-1} p_k(s_{k+1} | s_k, a_k) q_k(a_k | s_{0:k}, a_{0:k-1})$$

□

Theorem 2.2 - Markov Chains in Markov Decision Process

Let the policy π be Markovian. Then, $\{X_t\}_{t \in T}$ and $\{(X_t, Y_t)\}_{t \in T}$ are *Markov Chains*. Further, if the transition and decision probabilities are stationary (ie independent of t), then these chains are homogeneous.

NOTE - We only prove that the sequence of states $\{X_t\}_{t \in T}$ is a Markov Chain, see **Proof 2.2**.

Theorem 2.3 - Transition Kernel of Sequence of States $\{X_t\}_{t \in T}$

The sequence states $\{X_t\}_{t \in T}$ is a markov chain with transition kernel

$$\mathbb{P}(X_{t+1} = s' | X_t = s) = \sum_{a \in A} p_t(s' | s, a) q_t(a | s)$$

NOTE - This is proved in **Proof 2.2**

Proof 2.2 - Sequence of States $\{X_t\}_{t \in T}$ is a Markov Chain

To show that $\{X_t\}_{t \in T}$ is a markov chain it is sufficient to show

$$\mathbb{P}(X_{t+1} = s_{t+1} | X_{0:t} = s_{0:t}) = \mathbb{P}(X_{t+1} = s_{t+1} | X_t = s_t)$$

(ie is independent of $s_{0:t-1}$).

By the definition of conditional probabilities we have

$$\mathbb{P}(X_{t+1} = s_{t+1} | X_{0:t} = s_{0:t}) = \frac{\mathbb{P}(X_{0:t+1} = s_{0:t+1})}{\mathbb{P}(X_{0:t} = s_{0:t})}$$

By marginalising we can re-express the numerator as

$$\mathbb{P}(X_{0:t+1} = s_{0:t+1}) = \sum_{a_{0:t} \in A^{t+1}} \mathbb{P}(X_{0:t+1} = s_{0:t+1}, Y_{0:t} = a_{0:t})$$

Further, we can expand each term of this summation as

$$\begin{aligned}
& \frac{\mathbb{P}(X_{0:t+1} = s_{0:t+1}, Y_{0:t} = a_{0:t})}{\mathbb{P}(X_{0:t+1} = s_{0:t+1}, Y_{0:t} = a_{0:t})} \cdot \frac{\mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t} = a_{0:t})}{\mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1})} \cdot \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \\
&= \mathbb{P}(X_{t+1} = s_{t+1} | X_{0:t} = s_{0:t}, Y_{0:t} = a_{0:t}) \mathbb{P}(Y_t = a_t | X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \\
&=
\end{aligned}$$

Substituting this into the previous expression we get

$$\mathbb{P}(X_{0:t+1} = s_{0:t+1}) = \sum_{a_{0:t} \in A^{t+1}} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1})$$

We can make this expression recursive, by splitting the summation into two stages: one over $X_{0:t-1}$; the other over X_t .

$$\begin{aligned} & \mathbb{P}(X_{0:t+1} = s_{0:t+1}) \\ &= \sum_{a_{0:t} \in A^{t+1}} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \\ &= \sum_{a_t \in A} \sum_{a_{0:t-1} \in A^t} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \\ &= \left(\sum_{a_t \in A} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \right) \left(\sum_{a_{0:t-1} \in A^t} \mathbb{P}(X_{0:t} = s_{0:t}, Y_{0:t-1} = a_{0:t-1}) \right) \\ &= \left(\sum_{a_t \in A} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \right) \mathbb{P}(X_{0:t} = s_{0:t}) \text{ by the recursive definition} \end{aligned}$$

This expression can be substituted into the expression of the conditional probability

$$\begin{aligned} \mathbb{P}(X_{t+1} = s_{t+1} | X_{0:t} = s_{0:t}) &= \frac{\mathbb{P}(X_{0:t+1} = s_{0:t+1})}{\mathbb{P}(X_{0:t} = s_{0:t})} \\ &= \frac{(\sum_{a_t \in A} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t)) \mathbb{P}(X_{0:t} = s_{0:t})}{\mathbb{P}(X_{0:t} = s_{0:t})} \\ &= \sum_{a_t \in A} p_t(s_{t+1}|s_t, a_t) q_t(a_t|s_t) \end{aligned}$$

This final expression is independent of $s_{0:t-1}$, thus $\{X_t\}_{t \in T}$ is a markov chain. Further, it has transition kernel

$$\mathbb{P}(X_{t+1} = s' | X_t = s) = \sum_{a \in A} p_t(s'|s, a) q_t(a|s)$$

Remark 2.2 -

In *Markov Decision Processes* expressions of the following form

$$\mathbb{P}(X_0 = s_0, \dots, X_t = X_t, Y_0 = a_0, \dots, Y_t = a_t) \quad \mathbb{E}[f(X_0, \dots, X_t, Y_0, \dots, Y_t)]$$

depend on the applied policy π . To emphasis this the following notation is used $\mathbb{P}^\pi(\cdot)$ and $\mathbb{E}^\pi[\cdot]$.

$$\mathbb{P}^\pi(X_0 = s_0, \dots, X_t = X_t, Y_0 = a_0, \dots, Y_t = a_t) \quad \mathbb{E}^\pi[f(X_0, \dots, X_t, Y_0, \dots, Y_t)]$$

2.2 Markov Decision Processes

Definition 2.6 - Markov Decision Process (MDP)

A *Markov Decision Process* (MDP) is a *Sequential Decision Problem* where the underlying stochastic system evolves in a Markovian Fashion. There are two main components to a *Markov Decision Process*:

- i). A stochastic system.

- ii). An agent.

Proposition 2.5 - Process of an MDP

A *Markov Decision Process* follows the following steps

- i). At a specified point in time, an agent observes the state of a system.
- ii). Based on the observed state, the agent selects an action.
- iii). The selected action produces two results
 - An immediate reward for the agent.
 - A new evolved state for the system.
- iv). The process repeats with the updated system.

2.2.1 Problem Formulation

Definition 2.7 - Markov Decision Problem

In a *Markov Decision Problem* the agent is tasked with maximising the total reward recieved over a given time horizon T .

$$\text{TotalReward} := \sum_{t \in T} r_t(X_t, Y_t)$$

A *Markov Decision Problem* has the following elements

- i). *Decision Epochs* T - The points in time where the agent selects and makes their action.
- ii). *System States* S - A numerical representation of the condition of the stochastic system.
- iii). *Actions* A - The product of the agent's decision. Actions affect the behaviour of the system.
- iv). *Immediate Rewards* $r_t(s, a)$ - What the agents receives for their action. This is a measure of the quality of the action taken (More reward=Better Action).
- v). *Transition Probabilities* $p_t(\cdot|s, a)$ - The probability of the system being in a particular state at the next time step $t + 1$, given it's current state s and the action a the agent takes. These have *Markovian Properties*.

Definition 2.8 - Time-Horizon T

The *Time-Horizon* is the set T of all *Decision Epochs*. There are two types of *Time-Horizon*

- i). *Continuous-Time* where the time-horizon is specified over an interval $T := [t_1, t_2]$
- ii). *Discrete-Time* where the time-horizon is specified for specific points in time $T := \{t_0, \dots, t_N\}$ with $t_0 < \dots < t_N$. As the intervals between epochs is irrelevant T is often defined as a subset of natural numbers $T := \{0, \dots, N\}$. There are two further categories of *Discrete Time-Horizons*
 - *Finite Time* - $T := \{0, \dots, N\}$.
 - *Infinite Time* - $T := \{0, 1, \dots\}$.

N.B. - Only discrete-time time-horizons are within the scope of this course.

Definition 2.9 - State-Space S

The *State-Space* is the set S of all states the stochastic system can take. There are three types of *State Space*

- i). *Continuous-State* - The state space S is uncountable.
- ii). *Discrete-State* - The state space S is countably infinite.
- iii). *Finite-State* - The state space S is countably finite $S := \{s_1, \dots, s_N\}$ with $N \geq 2$.

N.B. - Only finite-state state-spaces are within the scope of this course.

Definition 2.10 - Action-Space A

The *Action-Space* is the set A of all actions available to the agent. There are three types of *Action Space*

- i). *Continuous-Action* - The action space A is uncountable.
- ii). *Discrete-Action* - The action space A is countably infinite.
- iii). *Finite-Action* - The action space A is countably finite $A := \{a_1, \dots, a_N\}$ with $N \geq 2$.

It is possible that not all actions in A are available for all states in S . We let $A(s)$ denote the set of available actions in state s .

N.B. - Only finite-action action-spaces are within the scope of this course.

Definition 2.11 - Transition Probabilities $p_t(\cdot|s, a)$

Transition Probabilities $p_t(\cdot|s, a)$ are a family of parametric probability mass functions on state-space S . Each transition probability is parameterised by the epoch t and are conditional on the system state s and action taken a .

The *Transition Probabilities* have *Markovian Properties*

$$\begin{aligned} \mathbb{P}(X_{t+1} = s_{t+1} | X_0 = s_0, Y_0 = a_0, \dots, X_t = s_t, Y_t = a_t) &= \mathbb{P}(X_{t+1} = s_{t+1} | X_t = s_t, Y_t = a_t) \\ &=: p_t(s_{t+1} | s_t, a_t) \end{aligned}$$

i.e. The probability of transition a given state s_{t+1} only depends on the current state s_t and action a_t (and nothing earlier).

$$\begin{aligned} \{X_{t+1} | X_0, Y_0, \dots, X_t, Y_t\} &\sim \{X_{t+1} | X_t, Y_t\} \\ &\sim p_t(\cdot | X_t, Y_t) \end{aligned}$$

Definition 2.12 - Decision Rules

A *Decision Rule* is a procedure for selecting an action at the specified decision epoch. In the process of selecting an action, the decision rule takes into account the current system state, past system state and past agent actions.

There are four types of *Decision Rule*

- i). *History Dependent Randomised (HR)* - The decision rule $q_t(\cdot)$ is a conditional probability mass function on the action-space A which represent the probability of taking a given action given all currently available information.

$$\begin{aligned} \mathbb{P}(Y_0 = a_0 | X_0 = s_0) &= q_0(a_0 | s_0) \\ \mathbb{P}(Y_t = a_t | X_0 = s_0, Y_0 = a_0, \dots, X_{t-1} = s_{t-1}, Y_{t-1} = a_{t-1}, X_t = s_t) &= q_t(a_t | s_0, \dots, s_t, a_0, \dots, a_{t-1}) \\ &\implies \{Y_0 | X_0\} \sim q_0(\cdot | X_0) \\ \{Y_t | X_0, \dots, X_t, Y_0, \dots, Y_{t-1}\} &\sim q_t(\cdot | X_0, \dots, X_t, Y_0, \dots, Y_{t-1}) \end{aligned}$$

- ii). *History Dependent Deterministic (HD)* - The decision rule $d_t(\cdot)$ is a deterministic function of all currently available information

$$Y_t := d_t(X_0, \dots, X_t, Y_0, \dots, Y_{t-1})$$

- iii). *Markovian Randomised (MR)* - The decision rule $q_t(\cdot)$ is a conditional probability mass function on the action-space A which is conditional only on the current state (ie is independent of past states and actions)

$$\begin{aligned} \mathbb{P}(Y_t = a_t | X_0 = s_0, Y_0 = a_0, \dots, X_{t-1} = s_{t-1}, Y_{t-1} = a_{t-1}, X_t = s_t) &= q_t(a_t | s_t) \\ \implies \{Y_t | X_0, \dots, X_t, Y_0, \dots, Y_{t-1}\} &\sim q_t(\cdot | X_t) \end{aligned}$$

- iv). *Markovian Deterministic (MD)* - The decision rule $d(t)$ is a deterministic function of the current state

$$Y_t = d_t(X_t)$$

Remark 2.3 - M

arkovian Decision rules are memoryless.

Definition 2.13 - Decision Policy π

A *Decision Policy* π is a procedure which specifies a decision rule for any decision epoch. A *Decision Policy* is a sequence of decision rules

$$\pi = \{q_t(\cdot | \cdot)\}_{t \in T}$$

where $q_t(\cdot | \cdot)$ is the decision probability at epoch t .

There are two classes of decision policies

- i). *Stationary* - The decision rule is the same for all decision epochs (ie $q_t(\cdot | \cdot)$ is independent of t).
- ii). *Non-Stationary* - The decision rule varies between epochs (ie $q_t(\cdot | \cdot)$ is dependent on t).

Remark 2.4 - Quantifying Total Reward

As **TotalReward** is a random quantity, it cannot be maximised directly. Thus we target finding the policy π which gives the maximum expected reward, given the rewards $\{r_t(\cdot)\}_{t \in T}$ and transition probabilities $\{p_t(\cdot | \cdot)\}_{t \in T}$

$$\mathbb{E}^\pi[\text{TotalReward}] := \mathbb{E}^\pi \left(\sum_{t \in T} r_t(X_t, Y_t) \right)$$

2.2.2 Finite Horizon Problems

Definition 2.14 - *Finite Horizon Markov Decision Problems*

In a *Finite Horizon MDP* the agent makes decisions at a finite number of epochs $N < \infty$.

Proposition 2.6 - *Elements of Finite Horizon MDPs*

A *Finite Horizon MDP* over N epochs has the following steps

- Time Horizon $T = \{0, \dots, N - 1\}$.
- Transition Probabilities $p_0(s'|s, a), \dots, p_{N-1}(s'|s, a)$.
- Rewards $r_0(s, a), \dots, r_{N-1}(s, a), r_N(s)$.

Proposition 2.7 - *Objective in a Finite Horizon MDP*

For a *Finite Horizon MDP* with transition probabilities $p_0(s'|s, a), \dots, p_{N-1}(s'|s, a)$ and the rewards $r_0(s, a), \dots, r_{N-1}(s, a), r_N(s)$, our objective is to find a *History Dependent Randomised* policy over time-horizon T $\pi \in HR(T)$ which achieves the maximum expected reward

$$R^\pi := \mathbb{E}^\pi \left[\left(\sum_{t=0}^{N-1} r_t(X_t, Y_t) \right) + r_N(X_N) \right]$$

Definition 2.15 - *Terminal Reward $r_X(s)$*

The *Terminal Reward* is recieved for the final state the agent ends up in (ie the state they are in after N actions). As no action is taken at epoch $t = N$, $r_N(\cdot)$ depends only on the state s .

Proposition 2.8 - *Stochastic Systems for a Finite Horizon MDP*

As the sequence of states $\{X_t\}_{t \in T}$ is a markov chain we can deduce the probability of transitioning to a given state, given all available information.

$$\forall t \in T \quad (X_{t+1}|X_{0:t}, Y_{0:t}) \sim^{[1]} (X_{t+1}|X_t, Y_t) \sim p_t(X_{t+1}|X_t, Y_t)$$

[1] by Markov property.

Two-Stage MDP

Definition 2.16 - *Two-Stage MDP*

A *Two-Stage MDP* is a *Finite Horizon MDP* over only $N = 2$ epochs. It has the following elements

- Time Horizon $T = \{0, 1\}$.
- Transition Probabilities $p_0(s'|s, a), p_1(s'|s, a)$.
- Rewards $r_0(s, a), r_1(s, a), r_2(s)$.

Proposition 2.9 - *Stochastic System for a Two-Stage MDP*

$$\begin{aligned} (X_2|X_{0:1}, Y_{0:1}) &\sim (X_2|X_1, Y_1) \sim p_1(\cdot|X_1, Y_1) \\ (X_1|X_0, Y_0) &\sim p_0(\cdot|X_0, Y_0) \end{aligned}$$

Proposition 2.10 - Objective in a Two-Stage MDP

For a *Two-Stage MDP* with transition probabilities $p_0(s'|s, a), p_1(s'|s, a)$ and the rewards $r_0(s, a), r_1(s, a), r_2(s)$, our objective is to find a *History Dependent Randomised* policy over time-horizon T $\pi \in HR(T)$ which achieves the maximum expected reward

$$\begin{aligned} R^\pi &:= \mathbb{E}^\pi [r_0(X_0, Y_0) + r_1(X_1, Y_1) + u_2^*(X_2)] \\ &= \mathbb{E}^\pi [r_0(X_0, Y_0) + r_1(X_1, Y_1)] + \mathbb{E}^\pi [u_2^*(X_2)] \end{aligned}$$

where $u_2^*(X_2) := r_2(s)$.

Proposition 2.11 - Optimising a Two-Stage MDP - Epoch $t = 1$

For a *Two-Stage MDP* we want to maximise the expected total reward R^π

$$R^\pi := \mathbb{E}^\pi [r_0(X_0, Y_0) + r_1(X_1, Y_1)] + \mathbb{E}^\pi [u_2^*(X_2)]$$

Note that by the *Tower Property*

$$\mathbb{E}^\pi [u_2^*(X_2)] = \mathbb{E}^\pi [\mathbb{E}^\pi (u_2^*(X_2) | X_1, Y_1)]$$

Thus,

$$\begin{aligned} R^\pi &= \mathbb{E}^\pi [r_0(X_0, Y_0) + r_1(X_1, Y_1)] + \mathbb{E}^\pi [\mathbb{E}^\pi (u_2^*(X_2) | X_1, Y_1)] \\ &= \mathbb{E}^\pi [r_0(X_0, Y_0) + r_1(X_1, Y_1) + \mathbb{E}^\pi (u_2^*(X_2) | X_1, Y_1)] \end{aligned}$$

By the definition of conditional expectation and transition probabilities we have

$$\begin{aligned} \mathbb{E}^\pi (u_2^*(X_2) | X_1, Y_1) &= \sum_{s' \in S} u_2^*(s') \mathbb{P}^\pi (X_2 = s' | X_1, Y_1) \\ &= \sum_{s' \in S} u_2^*(s') p_1(s' | X_1, Y_1) \end{aligned}$$

This can be substituted back into the expression of expected total reward to get an expression which is sufficient to choose an optimal action at epoch $t = 1$

$$R^\pi = \mathbb{E}^\pi \left[r_0(X_0, Y_0) + r_1(X_1, Y_1) + \sum_{s' \in S} u_2^*(s') p_1(s' | X_1, Y_1) \right]$$

At epoch $t = 1$ X_0, Y_0, X_1 are all known. So, the expression $r_0(X_0, Y_0) + r_1(X_1, Y_1) + \sum_{s' \in S} u_2^*(s') p_1(s' | X_1, Y_1)$ can be interpreted as a deterministic function of Y_1 (the action played at epoch $t = 1$)

$$\begin{aligned} Y_1^* &:= \operatorname{argmax}_{a \in A(s)} \left(r_0(X_0, Y_0) + r_1(X_1, a) + \sum_{s' \in S} u_2^*(s') p_1(s' | X_1, a) \right) \\ &= \operatorname{argmax}_{a \in A(s)} \left(r_1(X_1, a) + \sum_{s' \in S} u_2^*(s') p_1(s' | X_1, a) \right)^{[1]} \end{aligned}$$

I now define the following notation

$$\begin{aligned} u_1^*(s) &:= \max_{a \in A(s)} \left(r_1(s, a) + \sum_{s' \in S} u_2^*(s') p_1(s' | s, a) \right) \\ d_1^*(s) &:= \operatorname{argmax}_{a \in A(s)} \left(r_1(s, a) + \sum_{s' \in S} u_2^*(s') p_1(s' | s, a) \right) \end{aligned}$$

^[1]As terms independent of a do not affect the maximisation.

$u_1^*(s)$ is the maximum expected reward which can be achieved after being in state s at epoch $t = 1$, $d_1^*(s)$ is optimal decision function (ie optimal action) given the agent starts epoch $t = 0$ in state s (There may be multiple such actions). Thus

$$Y_1^* = d_1^*(X_1)$$

As $u_1^*(s)$ is the optimal value we can use it to bound the expected reward expression

$$\forall s \in S, \forall a \in A(s) \quad r_1(s, a) + \sum_{s' \in S} u_2^*(s') p_1(s'|s, a) \leq u_1^*(s)$$

This inequality becomes an equality when $a = d_1^*(s)$.

Consider where $s = X_1$ and $a = Y_1$, we can conclude

$$r_1(X_1, Y_1) + \sum_{s' \in S} u_2^*(s') p_1(s'|X_1, Y_1) \leq u_1^*(X_1)$$

This inequality becomes an equality when $Y_1 = d_1^*(X_1)$.

To consider the first action Y_0 we add $r_0(X_0, Y_0)$ to both sides and take expectations, giving

$$\mathbb{E}^\pi \left(r_0(X_0, Y_0) + r_1(X_1, Y_1) + \sum_{s' \in S} u_2^*(s') p_1(s'|X_1, Y_1) \right) \leq \mathbb{E}^\pi (r_0(X_0, Y_0) + u_1^*(X_1))$$

Again, this inequality becomes an equality when $Y_1 = d_1^*(X_1)$.

As the LHS is just the definition for total expected reward R^π we can restate the bound as a bound on R^π , nothing the inequality is maximised for $Y_1 = d_1^*(X_1)$

$$R^\pi \leq \mathbb{E}^\pi (r_0(X_0, Y_0) + u_1^*(X_1))$$

Proposition 2.12 - Optimising a Two-Stage MDP - Epoch $t = 0$

Let $\tilde{R}^\pi := \mathbb{E}^\pi (r_0(X_0, Y_0) + u_1^*(X_1))$ denote the bound derived in **Proposition 2.11**. From **Proposition 2.11** we have that $R^\pi \leq \tilde{R}^\pi$ and that this is an equality if $Y_1 = d_1^*(X_1)$.

By the *Tower Property*, we have

$$\mathbb{E}^\pi (U_1^*(X_1)) = \mathbb{E}^\pi [\mathbb{E}^\pi (u_1^*(X_1) | X_0, Y_0)]$$

Using this we can restate \tilde{R}^π as

$$\begin{aligned} \tilde{R}^\pi &= \mathbb{E}^\pi (r_0(X_0, Y_0)) + \mathbb{E}^\pi (u_1^*(X_1)) \\ &= \mathbb{E}^\pi (r_0(X_0, Y_0)) + \mathbb{E}^\pi [\mathbb{E}^\pi (u_1^*(X_1) | X_0, Y_0)] \end{aligned}$$

By the definition of conditional expectations and transition probabilities we have

$$\begin{aligned} \mathbb{E}^\pi (u_1^*(X_1) | X_0, Y_0) &= \sum_{s' \in S} u_1^*(s') \mathbb{P}^\pi (X_1 = s' | X_0, Y_0) \\ &= \sum_{s' \in S} u_1^*(s') p_0(s' | X_0, Y_0) \\ \implies \tilde{R}^\pi &= \mathbb{E}^\pi \left[r_0(X_0, Y_0) + \sum_{s' \in S} u_1^*(s') p_0(s' | X_0, Y_0) \right] \end{aligned}$$

At epoch $t = 0$ the value of X_0 is known, thus the following is a deterministic function of Y_0

$$r_0(X_0, Y_0) + \sum_{s' \in S} u_1^*(s') p_0(s' | X_0, Y_0)$$

This means the optimal action Y_0^* at epoch $t = 0$ is the action which maximises this function

$$Y_0^* := \operatorname{argmax}_{a \in A(s)} \left(r_0(X_0, a) + \sum_{s' \in S} u_1^*(s') p_0(s'|X_0, a) \right)$$

Define $u_0^*(s)$ to be the maximum reward after epoch $t = 0$ given the agent starts the epoch in state s and d_0^* be the optimal decision function (ie optimal action) at epoch $t = 0$ given the agent starts the epoch in state s .

$$\begin{aligned} u_0^*(s) &:= \max_{a \in A(s)} \left(r_0(s, a) + \sum_{s' \in S} u_1^*(s') p_0(s'|s, a) \right) \\ d_0^*(s) &:= \operatorname{argmax}_{a \in A(s)} \left(r_0(s, a) + \sum_{s' \in S} u_1^*(s') p_0(s'|s, a) \right) \end{aligned}$$

Using this $Y_0^* = d_0^*(X_0)$ and we can bound the expected reward

$$\forall s \in S, a \in A(s) \quad r_0(s, a) + \sum_{s' \in S} u_1^*(s') p_0(s'|s, a) \leq u_0^*(s)$$

Note that is becomes an equality if $a = d_0^*(s)$.

Setting $s = X_0, a = Y_0$ we can conclude

$$r_0(X_0, Y_0) + \sum_{s' \in S} u_1^*(s') p_0(s'|X_0, Y_0) \leq u_0^*(X_0)$$

This bound is an equality if $a = d_0^*(X_0)$. We can therefore bound the expected reward as

$$\begin{aligned} \mathbb{E}^\pi \left[r_0(X_0, Y_0) + \sum_{s' \in S} u_1^*(s') p_0(s'|X_0, Y_0) \right] &\leq \mathbb{E}^\pi [u_0^*(X_0)] \\ \Rightarrow \tilde{R}^\pi &\leq \mathbb{E}^\pi [u_0^*(X_0)] \end{aligned}$$

which is optimised when $Y = d_0^*(X_0)$.

As the LHS is the definition of \tilde{R}^π we can bound the total expected reward as

$$R^\pi \leq \tilde{R}^\pi \leq \mathbb{E}^\pi [u_0^*(X_0)]$$

which is optimised if $Y_0 = d_0^*(X_0)$ and $Y_1 = d_1^*(X_1)$.

Remark 2.5 - Conclusions

- The expected total reward R^π cannot exceed $\mathbb{E}^\pi [u_0^*(X_0)]$, but it can equal it.
- The optimal expected total reward is achieved when $Y_0 = d_0^*(X_0), Y_1 = d_1^*(X_1)$. Thus the optimal policy π^* is

$$\pi^* := (d_0^*(s), d_1^*(s))$$

Note that the construction of π^* is from functions $u_0^*(s), u_1^*(s), d_0^*(s), d_1^*(s)$ which are defined for $t \in \{0, 1\}$ as

$$\begin{aligned} d_t^* &= \operatorname{argmax}_{a \in A(s)} (r_t(s, a) + \sum_{s' \in S} u_{t+1}^*(s') p_t(s'|s, a)) \\ u_t^* &= \max_{a \in A(s)} (r_t(s, a) + \sum_{s' \in S} u_{t+1}^*(s') p_t(s'|s, a)) \end{aligned}$$

where $u_N^*(s) := r_N(s)$

Remark 2.6 - Optimal Policy for any Finite-Horizon MDP

We can generalise the optimal policy for a *Two-Stage MDP*, given in **Remark 2.5**, to get an optimal decision policy for any *Finite-Horizon MDP*.

For a *Finite-Horizon MDP* over N epochs the optimal decision policy is

$$\pi^* := (d_0^*(s), \dots, d_{N-1}^*(s))$$

where for $t \in \{0, \dots, N-1\}$

$$\begin{aligned} d_t^* &= \operatorname{argmax}_{a \in A(s)} (r_t(s, a) + \sum_{s' \in S} u_{t+1}^*(s') p_t(s'|s, a)) \\ u_t^* &= \max_{a \in A(s)} (r_t(s, a) + \sum_{s' \in S} u_{t+1}^*(s') p_t(s'|s, a)) \end{aligned}$$

General Finite-Horizon MDPs**Definition 2.17 - Optimality Equation $u_k^*(s)$**

The *Optimality Equation* for a General Finite-Horizon MDP over N epochs is defined by the following backward recursion

$$u_k^*(s) := \max_{a \in A(s)} \left(r_k(s, a) + \sum_{s' \in S} u_{k+1}^*(s') p_k(s'|s, a) \right) \text{ for } k \in [0, N-1], s \in S, u_N^*(\cdot) := r_N(\cdot)$$

AKA *Dynamic Programming Algorithm*

Definition 2.18 - Optimal Decision Rules $d_k^*(\cdot)$

The *Optimality Decision Rules* for a General Finite-Horizon MDP over N epochs is defined by the following backward recursion

$$d_k^*(s) := \operatorname{argmax}_{a \in A(s)} \left(r_k(s, a) + \sum_{s' \in S} u_{k+1}^*(s') p_k(s'|s, a) \right) \text{ for } k \in [0, N-1], s \in S$$

Definition 2.19 - Optimal Policy π^*

The *Optimal Policy* for a General Finite-Horizon MDP over N epochs is defined as

$$\pi^* := \{d_t^*(s)\}_{t \in T}$$

Theorem 2.4 - Bound on Expected Total Reward

For any *History Dependent Randomised Policy* $\pi \in HR(T)$ over time-horizon T the following inequality holds for the expected Total Reward

$$\underbrace{\mathbb{E}^\pi \left(\sum_{t=0}^{N-1} r_t(X_t, Y_t) + r_N(X_N) \right)}_{\text{Expected Total Reward}} \leq \mathbb{E}^\pi(u_0^*(X_0)) = \sum_{s \in S} u_0^*(s) \mathbb{P}(X_0 = s)$$

This inequality becomes an equality if the policy is optimal $\pi = \pi^*$.

Proof 2.3 - *Theorem 2.4*

Remark 2.7 - *The Initial State X_0 is not affect by any policy π*

This can be emphasized by dropping policy notation

$$\begin{aligned} \mathbb{P}(X_0 = s) &\text{ instead of } \mathbb{P}^\pi(X_0 = s) \\ \mathbb{E}(u_0^*(X_0)) &\text{ instead of } \mathbb{E}^\pi(u_0^*(X_0)) \end{aligned}$$

Definition 2.20 - *Value Function $v^\pi(\cdot)$*

The *Value Function* associated with a *History Dependent Randomised Policy* $\pi \in HR(T)$ is the expected total reward, given the initial state $s \in S$ (ie $X_0 = s$)

$$v^\pi(s) := \mathbb{E}^\pi \left(\sum_{t=0}^{N-1} r_t(X_t, Y_t) + r_N(X_N) \mid X_0 = s \right)$$

Definition 2.21 - *Optimal Value Function $v^*(\cdot)$*

The *Optimal Value Function* is the maximum expected total reward, given the initial state $s \in S$ (ie $X_0 = s$)

$$\begin{aligned} v^*(s) &:= \max_{\pi \in HR(T)} v^\pi(s) \\ &= \max_{\pi \in HR(T)} \mathbb{E}^\pi \left(\sum_{t=0}^{N-1} r_t(X_t, Y_t) + r_N(X_N) \mid X_0 = s \right) \end{aligned}$$

Note that this considers all possible *History Dependent Randomised Policies*.

Theorem 2.5 - *Restatements of Optimal Value Function*

Here are two equivalent expressions of the *Optimal Value Function* $v^*(\cdot)$

- i). $v^*(s) = u_0^*(s) \forall s \in S$
- ii). $v^*(s) = v^{\pi^*}(s) \forall s \in S$

Optimality Principle

Definition 2.22 - *Tail Subproblem*

The *Tail Subproblem of Length $l \in [1, N]$* is a subproblem of a *Finite Time-Horizon MDP* over N epochs. The subproblem has the following elements

- L decision epochs.
- Time-Horizon $T_L = \{N - L, \dots, N - 1\}$.
- Transition Probabilities $p_{N-L}(s'|s, a), \dots, p_{N-1}(s'|s, a)$.
- Rewards $r_{N-L}(s, a), \dots, r_{N-1}(s, a), r_N(s)$.

All these elements are from the main problem (ie this is a subproblem). Further, this subproblem is equivalent to the last L steps of the *Finite Time-Horizon MDP* over N epochs. The practical difference between the problem and the subproblem is that the initial epoch for the subproblem is $t = N - L$, rather than $t = 0$.

The object of this subproblem is to find the *History Dependent Randomised Policy* $\pi \in HR(T_L)$ which maximises the expected total reward

$$\operatorname{argmax}_{\pi \in HR(T_L)} \mathbb{E}^\pi \left[\left(\sum_{t=N-L}^{N-1} r_t(X_t, Y_t) \right) + r_N(X_N) \right]$$

Remark 2.8 - *Equivalence of MDP and Tail Subproblem*

If the tail subproblem is translated from epochs $N - L, \dots, N - 1$ to epochs $0, \dots, L - 1$ it is equivalent an MDP over L epochs (The elements of the problem need to translated accordingly).

This means the optimality equation for the tail subproblem is just a re-indexing of the optimality equations for the main problem.

Proposition 2.13 - *Optimality Equation for Tail Subproblem $u_{l,k}^*(\cdot)$*

The *Optimality Equation* for the tail subproblem of length L is defined with the following backwards recursion

$$u_{L,k}^*(s) := \max_{a \in A(s)} \left(r_k(s, a) + \sum_{s' \in S} u_{l,k+1}^*(s') p_k(s'|s, a) \right)$$

for $k \in [N - L, N - 1], s \in S$ and with $u_{l,N}^*(\cdot) := r_N(\cdot)$.

Proposition 2.14 - *Optimality Decision Rules for Tail Subproblem $d_{l,k}^*(\cdot)$*

The *Optimality Decision Rules* for the tail subproblem of length L is defined with the following backwards recursion

$$d_{L,k}^*(s) := \operatorname{argmax}_{a \in A(s)} \left(r_k(s, a) + \sum_{s' \in S} u_{l,k+1}^*(s') p_k(s'|s, a) \right)$$

for $k \in [N - L, N - 1], s \in S$.

Proposition 2.15 - *Optimality Decision Policy for Tail Subproblem π_l^**

The *Optimality Decision Policy* for the tail subproblem of length L is defined

$$\pi_l^* = \{d_{l,t}^*(s)\}_{t \in T_L}$$

Remark 2.9 - *Properties of Optimality Equations, Decision Rules and Policy*

- The optimality equation for the tail subproblem $u_{l,k}^*(\cdot)$ is a subrecursion of the optimality equation for the main problem $u_k^*(\cdot)$.
- The initial condition in the optimality equation for the tail subproblem $u_{l,N}^*(\cdot) = r_N(\cdot)$ is the same as the initial condition for the main problem $u_N^*(\cdot) = r_N(\cdot)$.

Given these two properties, the equations generated by the optimality equation for the tail subproblem are identical to the corresponding functions generated by the optimality equation for the main problem.

Theorem 2.6 - Optimality Equations for Sub & Main Problem are Identical

Consider a tail subproblem of length L , then $\forall s \in S, k \in [N - L, N - 1]$

$$\begin{aligned} u_{l,k}^*(s) &= u_k^*(s) \\ d_{l,k}^*(s) &= d_k^*(s) \end{aligned}$$

Definition 2.23 - Optimal Value Function $v_l^*(\cdot)$

The *Optimal value Function* for the tail subproblem of length L is defined as

$$v_l^*(s) := \max_{\pi \in HR(T_L)} \mathbb{E}^\pi \left[\sum_{t=N-L}^{N-1} r_t(X_t, Y_t) + r_N(X_N) \mid X_{N-L} = s \right] \quad \text{with } s \in S$$

The value of $v_l^*(s)$ can be interpreted as the maximum expected total reward recieved from epochs $N - L$ to N when the system is in state s at the start of epoch $t = N - L$.

Theorem 2.7 - Optimal Value Function as Optimality Equation

Consider a tail subproblem of length L , then $\forall s \in S, L \in [1, N]$ the following holds

$$v_L^*(s) = u_{L, N-L}^*(s)$$

Theorem 2.8 - Optimality Principle

Consider a tail subproblem of length L , then the *Optimality Principle States* that $\forall s \in S, L \in [1, N]$ the following hold

$$\begin{aligned} v_L^*(s) &= u_{N-L}^*(s) \\ \pi_L^* &= \{d_t^*(s)\}_{t \in T_L} \end{aligned}$$

Remark 2.10 - Optimality Principle

Here are some remarks about the *Optimality Principle*

- The *Optimality Principle* is a straightforward consequence of **Theorems 2.6 & 2.7**. Due to i), the *Dynamic Programming Algorithm*, starting with $L = N$ and going backwards towards $L = 0$, solves the tail suboptimal problem of length L .
- Due to ii), the *Optimal Policy* for the tail subproblem of length L π_L^* can be interpreted as the truncation of optimal policy for the main problem π^* to time horizon T_L .
- The *Dynamic Programming Algorithm* can be written as the following with $k \in [1, N], s \in S$ and $v_0^*(\cdot) := r_N(\cdot)$

$$\begin{aligned} v_k^*(s) &= \max_{a \in A(s)} \left(r_{N-k}(s, a) + \sum_{s' \in S} v_{k-1}^*(s') p_{N-k}(s' | s, a) \right) \\ d_k^*(s) &= \operatorname{argmax}_{a \in A(s)} \left(r_{N-k}(s, a) + \sum_{s' \in S} v_{k-1}^*(s') p_{N-k}(s' | s, a) \right) \end{aligned}$$

3 Probability

Definition 3.1 - Random Process

A *Random Process* is a collection of random variables indexed by time $\{X_t\}_{t \in T}$ (e.g. flipping a coin several times). Each of these random variables can take a value from a state space S . A random process is a *Discrete Time Process* if the index set T is discrete. A random process is a *Continuous Time Process* if the index set T is continuous.

3.1 Probability Inequalities

Remark 3.1 - We can use the moments of a random variable to determine bounds on the probability of it taking values in a certain set.

Theorem 3.1 - Markov's Inequality

Let X be a non-negative random variable. Then

$$\forall c > 0 \quad \mathbb{P}(X \geq c) \leq \frac{\mathbb{E}(X)}{c}$$

Proof

Consider an event A and define its indicator $\mathbb{1}(A)(\omega) := \begin{cases} 1 & \omega \in A \\ 0 & \omega \notin A \end{cases}$. Fix $c > 0$, then

$$\begin{aligned} \mathbb{E}(X) &\geq \mathbb{E}[X \mathbb{1}(X \geq c)] \\ &\geq \mathbb{E}[c \mathbb{1}(X \geq c)] \\ &= c \mathbb{P}(X \geq c) \\ \implies \mathbb{P}(X \geq c) &\leq \frac{1}{c} \mathbb{E}(X) \end{aligned}$$

Theorem 3.2 - Chebyshev's Inequality

Let X be a random-variable with finite mean and variance. Then

$$\forall c > 0 \quad \mathbb{P}(|X - \mathbb{E}(X)| \geq c) \leq \frac{\text{Var}(X)}{c^2}$$

Proof

Note that the events $|X - \mathbb{E}(X)| \geq c$ and $(X - \mathbb{E}(X))^2 \geq c^2$ are equivalent. Note that $\text{Var}([X - \mathbb{E}(X)]^2) = \text{Var}(X)$. Then the result follows by *Markov's Inequality*.

Theorem 3.3 - Chebyshev's Inequality for Sum of IIDs

Let X_1, \dots, X_n be IID random variables with finite mean μ and finite variance σ^2 .

$$\forall c > 0 \quad \mathbb{P}\left(\left|\sum_{i=1}^n X_i - n\mu\right| \geq nc\right) \leq \frac{\sigma^2}{nc^2}$$

Proof

This is proved by extending the proof of **Theorem 2.2** and noting that the variance of a sum of IIDs is the sum of the individual variances.

Theorem 3.4 - Chernoff Bounds

Let X be a random variable whose moment-generating function $\mathbb{E}[e^{\theta X}]$ is finite $\forall \theta$. Then

$$\forall c \in \mathbb{R} \quad \mathbb{P}(X \geq c) \leq \inf_{\theta > 0} e^{-\theta c} \mathbb{E}(e^{\theta X}) \quad \text{and} \quad \mathbb{P}(X \leq c) \leq \inf_{\theta < 0} e^{-\theta c} \mathbb{E}(e^{\theta X})$$

Proof

Note that the events $X \geq c$ and $e^{\theta X} \geq e^{\theta c}$ are equivalent for all $\theta > 0$. The result follows by applying *Markov's Inequality* to $e^{\theta X}$ and taking the best bound over all possible θ .

$$\begin{aligned}\mathbb{P}(X \geq c) &= \mathbb{P}(e^{\theta X} \geq e^{\theta c}) \\ &\leq e^{-\theta c} \mathbb{E}(e^{\theta X}) \\ &\leq \inf_{\theta > 0} e^{-\theta c} \mathbb{E}(e^{\theta X})\end{aligned}$$

Theorem 3.5 - Chernoff Bounds for Sum of IIDs

Let X_1, \dots, X_n be IID random variables. Then $\forall c \in \mathbb{R}$

$$\begin{aligned}\mathbb{P}\left(\sum_{i=1}^n X_i \geq nc\right) &\leq \inf_{\theta > 0} e^{-n\theta c} (\mathbb{E}[e^{\theta X}])^n \\ \mathbb{P}\left(\sum_{i=1}^n X_i \leq nc\right) &\leq \inf_{\theta < 0} e^{-n\theta c} (\mathbb{E}[e^{\theta X}])^n\end{aligned}$$

Theorem 3.6 - Jensen's Inequality

Let f be a *Convex Function* and X be a random variable. Then

$$\mathbb{E}[f(X)] \geq f(\mathbb{E}[X])$$

Theorem 3.7 - Bound on Moment Generating Function

Let X be a random variable taking values in $[0, 1]$ with finite expected value μ . Then we can bound the MGF of the centred random variable with

$$\forall \theta \in \mathbb{R} \quad \mathbb{E}[e^{\theta(X-\mu)}] \leq e^{\theta^2/8}$$

Proof (of weaker version)

Let X_1 be an independent copy of X , so both have mean μ . We can easily verify that $f(x) = e^{\theta x}$ is a convex function for all $\theta \in \mathbb{R}$. By *Jensen's Inequality* to $f(\cdot)$ and X_1

$$\mathbb{E}[e^{-\theta X_1}] \geq e^{-\theta \mathbb{E}[X_1]} = e^{-\theta \mu} \quad (1)$$

Consequently

$$\begin{aligned}\mathbb{E}[e^{\theta(X-X_1)}] &= \mathbb{E}[e^{\theta X}] \cdot \mathbb{E}[e^{-\theta X_1}] && \text{by independence} \\ &\geq \mathbb{E}[e^{\theta X}] \cdot e^{-\theta \mu} && \text{by (1)} \\ &= \mathbb{E}[e^{\theta(X-\mu)}] \\ \implies \mathbb{E}[e^{\theta(X-X_1)}] &\geq \mathbb{E}[e^{\theta(X-\mu)}]\end{aligned}$$

Since $X, X_1 \in [0, 1]$ then $(X - X_1) \in [-1, 1]$. As X, X_1 have the same distribution $\mathbb{E}(X - X_1) = 0$ and the distribution is symmetric around the mean.

Define random variable S which is independent of X, X_1 and takes values $\{-1, 1\}$, each with probability $p = \frac{1}{2}$. $S(X - X_1)$ has the same distribution as $(X - X_1)$ due to independence of S and symmetry of $(X - X_1)$. Hence

$$\begin{aligned}\mathbb{E}[e^{\theta(X-X_1)}] &= \mathbb{E}[e^{\theta S(X-X_1)}] && \text{by identical distribution} \\ &\leq \mathbb{E}[e^{\theta S}] && (2) \text{ since } (X - X_1) \in [-1, 1] \\ &= \frac{1}{2}(e^{\theta} + e^{-\theta}) && \text{by def. of expectation} \\ \implies \mathbb{E}[e^{\theta(X-X_1)}] &\leq \frac{1}{2}(e^{\theta} + e^{-\theta})\end{aligned}$$

Note that $f(x) = e^x + e^{-x}$ is increasing for $x \in (0, \infty)$; decreasing for $x \in (-\infty, 0)$; and symmetric around 0.

Using a *Taylor Series* we can observe that

$$\begin{aligned} \frac{1}{2}(e^\theta - e^{-\theta}) &= \sum_{n=0}^{\infty} \frac{\theta^{2n}}{(2n)!} && \text{by Taylor expansion of } e^x \\ &\leq \sum_{n=0}^{\infty} \frac{(\theta^2/2)^n}{n!} \\ &\stackrel{\text{def.}}{=} e^{\theta^2/2} \\ \implies \frac{1}{2}(e^\theta + e^{-\theta}) &\leq e^{\theta^2/2} \end{aligned}$$

Combining all these results we get

$$\begin{aligned} \mathbb{E}[e^{\theta(X-\mu)}] &\leq \mathbb{E}[e^{\theta(X-X_1)}] \leq \frac{1}{2}(e^\theta + e^{-\theta}) \leq e^{\theta^2/2} \\ \implies \mathbb{E}[e^{\theta(X-\mu)}] &\leq e^{\theta^2/2} \end{aligned}$$

□

Theorem 3.8 - Hoeffding's Theorem

Let X_1, \dots, X_n be IID random variables taking values in $[0, 1]$ and with finite expected value μ . Then

$$\forall t > 0 \quad \mathbb{P}\left(\sum_{i=1}^n (X_i - \mu) > nt\right) \leq e^{-2nt^2}$$

Proof

From *Chernoff's Bound* we have that

$$\forall \theta > 0 \quad \mathbb{P}\left(\sum_{i=1}^n (X_i - \mu) > nt\right) \leq e^{-\theta nt} \left(\mathbb{E}[e^{\theta(X-\mu)}]\right)^n$$

Using **Theorem 2.7** to bound the moment generating function, we get

$$\forall \theta > 0 \quad \mathbb{P}\left(\sum_{i=1}^n (X_i - \mu) > nt\right) \leq e^{-\theta nt} \cdot e^{n\frac{\theta^2}{8}} = e^{n(-\theta t + \frac{1}{8}\theta^2)}$$

Thus, by taking logs and rearranging, we get

$$\forall \theta > 0 \quad \frac{1}{n} \log \mathbb{P}\left(\sum_{i=1}^n (X_i - \mu) > nt\right) \leq -\theta t + \frac{\theta^2}{8}$$

We have that $-\theta t + \frac{\theta^2}{8}$ is minimised at $\theta = 4t$ which is positive if t is positive. Thus, by applying this bound and substituting $\theta = 4t$ we get

$$\forall \theta > 0 \quad \mathbb{P}\left(\sum_{i=1}^n (X_i - \mu) > nt\right) \leq e^{n(-4t^2 + \frac{1}{8}(16t^2))} = e^{n(-4t^2 + 2t^2)} = e^{-2nt^2}$$

□

3.2 Markov Processes

Definition 3.2 - Markov Property

A random process has the *Markov Property* if the conditional probability of a future state only depends on the current state.

$$\mathbb{P}(X_{t+1} = y | X_t = x_t, X_{t-1} = x_{t-1}) = \mathbb{P}(X_{t+1} = y | X_t = x_t)$$

A random process with the *Markov Property* is called a *Markov Process/Chain*.

Remark 3.2 - *On this course we only deal with discrete time markov chains*

Definition 3.3 - Transience

A state $x \in S$ is *Transient* if $\mathbb{P}(\exists t > 0 : X_t = x | X_0 = x) < 1$. The number of times the markov chain returns to a transient state is finite, with probability 1.

Definition 3.4 - Recurrent

A state $x \in S$ is *Recurrent* if $\mathbb{P}(\exists t > 0 : X_t = x | X_0 = x) = 1$. The number of times the markov chain returns to a recurrent state is infinite, with probability 1.

Every markov chain, with a finite state space S , has a recurrent communicating class.

Definition 3.5 - Communication Class

We say $y \in S$ is *Accessible* from $x \in S$ if $\exists t \geq 0$ st $[P^t]_{xy} > 0$.

We say x and y *communicate* (denoted xCy) if: x is *accessible* from y and y is *accessible* from x .

Communication is an *equivalence relation* on the state space S . Hence, *communication* partitions S into equivalence classes called *Communication Classes*. All elements of a *Communication Class* communicate with all other elements in the class, it is possible for elements to be accessible from another class but not for those elements to *communicate*.

If one state in a *Communicating Class* is *Transient/Recurrent* then all states are in that class.

If a *Markov Chain* has only one communicating class it is called *Irreducible*.

3.2.1 Discrete Time Markov Chains

Proposition 3.1 - Characterising a Discrete Time Markov Process

A *Discrete Time Markov Process* can be characterised by the set of all 1-step conditional probabilities

$$\mathbb{P}(X_{t+1} = y | X_t = x) \quad \forall x, y \in S$$

A markov chain is *time-homogeneous* if the 1-step conditionals only depend on x, y and not on t ($\mathbb{P}(X_{t+1} = y | X_t = x) = \mathbb{P}(X_1 = y | X_0 = x)$). The 1-step conditional probabilities of a *time-homogeneous markov process* can be specified in an $|S| \times |S|$ matrix P where

$$p_{x,y} = \mathbb{P}(X_{t+1} = y | X_t = x)$$

P is a *Stochastic Matrix*.

Proposition 3.2 - n-Step Transition Probabilities from 1-Step Transition Matrix

Let P be the 1-step transition matrix for a *time-homogeneous*.

The 2-step transition probabilities (ie $\mathbb{P}(X_{t+2} = z|X_t = x)$) can be found as

$$\begin{aligned}
 \mathbb{P}(X_{t+2} = z|X_t = x) &= \mathbb{P}(X_2 = z|X_0 = x) && \text{by time-homogeneity} \\
 &= \sum_{y \in S} \mathbb{P}(X_2 = z, X_1 = y|X_0 = x) \\
 &= \sum_{y \in S} \mathbb{P}(X_1 = y|X_0 = x) \mathbb{P}(X_2 = z|X_1 = y, X_0 = x) \\
 &= \sum_{y \in S} \mathbb{P}(X_1 = y|X_0 = x) \mathbb{P}(X_2 = z|X_1 = y) \\
 &= \sum_{y \in S} p_{xy} p_{yz} \\
 &\equiv [P^2]_{xz}
 \end{aligned}$$

This can be generalise for the n -step transition probabilities with

$$\mathbb{P}(X_{t+n} = z|X_t = x) = [P^n]_{xz}$$


Proposition 3.3 - Any Joint Probability from 1-Step Transition Matrix

For a time homogeneous markov process the joint distribution of any transition can be computed by considering the individual steps of the transition.

$$\mathbb{P}(X_{n_0} = x_0, X_{n_1} = x_1, X_{n_2} = x_2, \dots) = \mathbb{P}(X_{n_0} = x_0) \cdot [P^{n_1 - n_0}]_{x_0 x_1} \cdot [P^{n_2 - n_1}]_{x_1 x_2} \dots$$

Proposition 3.4 - State Diagram Representation

A graph/automata can be drawn to represent the transition probability matrix P . A node is assigned for each member of the state space S and an arrow is drawn between each pair of nodes (x, y) where $P_{xy} \neq 0$. Generally the value of P_{xy} is denoted on the arrow.



Definition 3.6 - Invariant Distribution

Let $\mu(t)$ denote the probability distribution of random variable X (i.e, $\mu_x(t) = \mathbb{P}(X_t = x)$). Then

$$\begin{aligned}
 \mu(t+1) &= \mathbb{P}(X_{t+1} = y) = \sum_{x \in S} \mathbb{P}(X_t = x, X_{t+1} = y) = \sum_{x \in S} \mu_x(t) p_{xy} = \mu(t)P \\
 \Rightarrow \mu(t+1) &= \mu(t)P
 \end{aligned}$$

A distribution π on the state space is called an *Invariant Distribution* if $\pi = \pi P$. If X_t has distribution π so will X_{t+1}, \dots . Every markov chain with a *finite* state space S has an *invariant distribution*. (Not necessarily true if S is infinite).

Proposition 3.5 - Finding an Invariant Distribution

If an *Invariant Distribution* it is easy to find by solving $\pi P = \pi$ and using normalising constant $\sum_{x \in S} \pi_x = 1$.

Remark 3.3 - If a Markov chain is irreducible, its invariant distribution (if one exists) is unique

If a Markov Chain is irreducible and has a finite state space, then it has a unique invariant distribution.

Example 3.1 - Markov Chains

- The *Asymmetric Simple Random Walk* on \mathbb{Z} is irreducible, transient and has no invariant distribution.
obvious not obvious
- The *Symmetric Simple Random Walk* on \mathbb{Z} is irreducible, recurrent and has no invariant distribution.
obvious not obvious

Theorem 3.9 - Ergodic Theorem for Markov Chains

Let $\{X_t\}_{t \in \mathbb{N}}$ be an irreducible markov chain on state space S (not necessarily finite) with unique invariant distribution π . Then

$$\forall x \in S \quad \lim_{t \rightarrow \infty} \frac{1}{t} \sum_{s=1}^t \mathbb{1}(X_s = x) = \pi_x$$

i.e. The fraction of time spend in state $x \in S$ tends to π_x in the long run.

Definition 3.7 - Period

The *Period* of a state $x \in S$ is the greatest common divisor of all possible return times to x

$$\text{Period}(x) := \gcd(\{t > 0 : \mathbb{P}(X_t = x | X_0 = x) > 0\})$$

A state $x \in S$ is *Aperiodic* if $\text{Period}(x) = 1$. An irreducible markov chain is *aperiodic* if all its states are aperiodic.

All states in a *communicating class* have the same period.

Proposition 3.6 - Marginal Distribution of Irreducible, Aperiodic Markov Chain

If an irreducible, aperiodic Markov Chain has an invariant distribution π , then

$$\forall x \in S \quad \mu_x(t) \xrightarrow{t \rightarrow \infty} \pi_x$$

Definition 3.8 - Reversibility

A markov chain $\{X_t\}_{t \in \mathbb{Z}}$ is *Reversible* if all joint distributions are the same forwards and backwards in time. (i.e. the distribution of the chain is the same if it was reversed).

An irreducible markov chain $\{X_t\}_{t \in \mathbb{Z}}$ with transition matrix P is reversible iff

$$\exists \pi \quad \text{st} \quad \pi_x p_{xy} = \pi_y p_{yx} \quad \forall x, y \in S$$

This is the *Local/Detailed Balance Equation*. Note that this is a system of $\binom{|S|}{2}$ equations which need to be consistent for reversibility to exist.

3.2.2 Continuous Time Markov Process

Definition 3.9 - Continuous Time Markov Process

A stochastic process $\{X_t\}_{t \in \mathbb{R}}$ is a *Continuous Time Markov Process* on state space S if

$$\forall s < t \ \& \ x, y \in S \quad \mathbb{P}(X_t = y | X_s = x, X_u, u \leq s) = \mathbb{P}(X_t = y | X_s = x)$$

ie future values only depend on the present value and not past.

If $\forall t, s, x, y \ \mathbb{P}(X_t = y | X_s = x)$ depends only on $x, y, t - s$ (observed values & change in time) then the process is *Time-Homogeneous*.

For *Time-Homogeneous Markov Processes* we let $P(t)$ denote the stochastic matrix with the probability of each possible transition after t time $[P(t)]_{xy} = \mathbb{P}(X_t = y | X_0 = x)$.

Remark 3.4 - A *Time-Homogeneous Markov Process* is completely described by its initial condition and the family of transition probability matrices $\{P(t) : t \geq 0\}$

This set of matrices $\{P(t) : t \geq 0\}$ is uncountably large.

Definition 3.10 - Chapman-Kolmogorov Equations

For a *Time-Homogeneous Markov Process* the family of stochastic matrices $\{P(t) : t \geq 0\}$ satisfy the following:

- i). $P(0) = I$;
- ii). $P(t + s) = P(t)P(s) = P(s)P(t)$

Hence

$$\frac{d}{dt}P(t) := \lim_{\delta \rightarrow 0} \frac{P(t + \delta) - P(t)}{\delta} = \lim_{\delta \rightarrow 0} \frac{\overbrace{P(t)P(\delta) - P(t)}^{\text{ii)}}}{\delta} = P(t) \lim_{\delta \rightarrow 0} \frac{(P(\delta) - I)}{\delta}$$

Suppose that $Q := \lim_{\delta \rightarrow 0} \frac{P(\delta) - P(0)}{\delta} = \lim_{\delta \rightarrow 0} \frac{P(\delta) - \overbrace{I}^{\text{i)}}}{\delta}$ exists.

Then $P(t)$ solve the following differential equations, known as the *Chapman-Kolmogorov Equations*

$$\frac{d}{dt}P(t) = \underbrace{P(t)Q}_{\text{forward eqn.}} = \underbrace{QP(t)}_{\text{backward eqn.}}$$

The solution to these equations is

$$P(t) = P(0)e^{Qt} = e^{Qt}P(0) = e^{Qt}I = e^{Qt}$$

N.B. Q is called the *Rate Matrix* or *Infinitesimal Generator* of the markov process.

Proposition 3.7 - Properties of the Rate Matrix, Q

Let Q be the rate matrix of a *continuous-time markov process*. Q has the following properties

- If $n \neq y$ then $q_{xy} := \lim_{\delta \rightarrow 0} \frac{[P(\delta)]_{xy} - 0}{\delta} \geq 0$. (The off-diagonal elements are non-negative).
- $\forall x \in S, \sum_{y \in S} q_{xy} = \lim_{\delta \rightarrow 0} \frac{1 - 1}{\delta} = 0$. The rows of Q sum to 0.

- Thus, the diagonal entries q_{xx} are negative. (We denote $-q_{xx}$ by q_x)

Proposition 3.8 - Interpreting the Rate Matrix Q

Let Q be the rate matrix of a *continuous-time markov process*.

If the markov process enters state x at time t , it will remain in x for a random time which is distributed $\text{Exp}(q_x)$. (Note that $q_x := -q_{xx}$).

It the jumps to state y with probability $\frac{q_{xy}}{q_x}$, independent of the past.

Definition 3.11 - Invariant Distributions

Suppose a *Continuous-Time Markov Process* starts with distribution $\boldsymbol{\mu}(0)$ on state space S (i.e. $\mathbb{P}(X_0 = x) = [\boldsymbol{\mu}(0)]_x$). Then, the distribution of X_t is $\boldsymbol{\mu}(t) := \boldsymbol{\mu}(0)P(t) = \boldsymbol{\mu}(0) \underbrace{e^{Qt}}_{\text{CK Eqns}}$.

If there exists a distribution $\boldsymbol{\pi}$ on the state space S st $\forall t \geq 0 \boldsymbol{\pi} = \boldsymbol{\pi}P(t) = \boldsymbol{\pi} \underbrace{e^{Qt}}_{\text{CK Eqns}}$, then $\boldsymbol{\pi}$ is an *Invariant Distribution*. This distribution is invariant wrt time.

If a markov process has a finite state space then it definitely has an invariant distribution.

Invariant Distributions are not guaranteed to be unique.

Proposition 3.9 - Finding an Invariant Distribution

Starting with $\boldsymbol{\pi} = \boldsymbol{\pi}e^{Qt}$ we find that differentiating wrt t and then evaluating at time $t = 0$ we get $\mathbf{0} = \boldsymbol{\pi}Q$ (The Global Balance Equations). This system of equations can be solved to find an *Invariant Distribution*.

A markov process is *reversible* iff there exists a distribution $\boldsymbol{\pi}$ on S which satisfies $\pi_x q_{xy} = \pi_y q_{yx} \forall x, y \in S$ (Local balance equations). Solving this system of equations will also find an *Invariant Distribution* but it is not guaranteed to have a solution.

Theorem 3.10 - Ergodic Theorem

Let $[X_t]_{t \in \mathbb{R}^+}$ is an *Irreducible Markov Process* on a state space S and has an invariant distribution $\boldsymbol{\pi}$. Then

$$\forall x \in S \quad \pi_x = \lim_{t \rightarrow \infty} \frac{1}{t} \int_0^t \mathbb{1}(X_s = x) ds$$

Moreover, for an arbitrary initial distribution $\boldsymbol{\mu}(0)$, $\boldsymbol{\mu}(t)$ converges to $\boldsymbol{\pi}$ pointwise (i.e. $\mu_x(t) \xrightarrow{t \rightarrow \infty} \pi_x$)

3.2.3 Poisson Process

Definition 3.12 - Counting Process

A *Counting Process* is a stochastic process $\{N(t)\}_{t \in \mathbb{R}}$ st

- $N(0) = 0$ and $N(t) \in \mathbb{Z}$ for all $t \geq 0$; and,
- $N(t)$ is a non-decreasing function of t .

Definition 3.13 - Independent Increments

A *Process* $\{N(t)\}_{t \in \mathbb{R}^+}$ is said to have *Independent Increments* if $\forall s \in (0, t)$, $(N(t) - N(s))$ is independent of $\{N(u) : u \in [0, s]\}$.

Definition 3.14 - Poisson Process

A *Poisson Process* is a *counting process* $\{N(t)\}_{t \in \mathbb{R}^+}$ which has independent increments and at least one of the following equivalence statements are true

- $\forall t \in [0, t] \quad (N(t) - N(s)) \sim \text{Po}(\lambda(t - s))$.
- $\mathbb{P}(N(t + \delta) - N(t) = 1) = \lambda\delta + o(\delta)$ and
 $\mathbb{P}(N(t + \delta) - N(t) = 0) = 1 - \lambda\delta + o(\delta)$ and
 $\mathbb{P}(N(t + \delta) - N(t) \geq 2) = o(\delta)$.
- The times between successive increments of the process $N(\cdot)$ are iid $\text{Exp}(\lambda)$ random variables.

The parameter $\lambda \in \mathbb{R}^{>0}$ is called the *rate* of the poisson process. *Poisson Processes* are continuous time markov chains.

Example 3.2 - Poisson Process

Counting the number of cars which have passed a given point over time.

Proposition 3.10 - Properties of Poisson Processes

Define $\{N(t)\}_{t \in \mathbb{R}^+}$ to be a *Poisson Process* with rate λ . Then the following properties hold

- The counting process $\{N(\beta t)\}_{t \in \mathbb{R}^+}$, with $\beta > 0$, is a Poisson Process with rate $\beta\lambda$.
- If $\{N_1(t)\}_{t \in \mathbb{R}^+}$ and $\{N_2(t)\}_{t \in \mathbb{R}^+}$ are independent poisson processes with rates λ_1 and λ_2 , respectively, then $\{N(t) := N_1(t) + N_2(t)\}_{t \in \mathbb{R}^+}$ is a poisson process with rate $\lambda := \lambda_1 + \lambda_2$.
- Let X_1, X_2, \dots be a sequence of iid $\text{Bern}(p)$ random variables, independent of $N(\cdot)$. Define $N_1(t) := \sum_{i=1}^{N(t)} X_i$ and $N_2(t) := \sum_{i=1}^{N(t)} (1 - X_i)$ (These are called *Bernoulli Thinnings*). These assign increments in $N(\cdot)$ randomly to either N_1 or N_2 (with probability p). Then, $N_1(\cdot)$ and $N_2(\cdot)$ are independent poisson processes with rates λp and $\lambda(1 - p)$, respectively.

3.3 Transformation of Random Variables

Example 3.3 - Discrete Case

Consider rolling a fair die where $\Omega := \{1, 2, 3, 4, 5, 6\}$ and $\forall \omega \in \Omega \quad \mathbb{P}(\omega) = \frac{1}{6}$.

Let $X(\omega) = \omega \quad \forall \omega \in \Omega$ so the pmf of X is given by

$$p_X(i) = \frac{1}{6} \quad \forall i \in \{1, \dots, 6\}$$

Consider $Y := X^2$. The pmf of Y is straightforward to work out as each value of X maps to a unique value of Y

$$P_Y(i) = \frac{1}{6} \quad \text{for } \sqrt{i} \in \{1, \dots, 6\}$$

Consider $Z := (X - 2)^2$. This is not quite so simple as multiple values of X map to the same value of Z .

$$p_Z(1) = \frac{2}{6}; \quad p_Z(i) = \frac{1}{6} \quad \text{for } i \in \{0, 4, 9, 16\}$$

Example 3.4 - Continuous Case

Let $X \sim \text{Uniform}[0, 1]$ and define $Y := 2X$. We have

$$f_X(x) = \begin{cases} 1 & x \in [0, 1] \\ 0 & \text{otherwise} \end{cases} \quad F_X(x) = \begin{cases} 0 & x < 0 \\ x & x \in [0, 1] \\ 1 & x > 1 \end{cases}$$

Now consider Y the cdf is

$$F_Y(y) := \mathbb{P}(Y \leq y) = \mathbb{P}(2X \leq y) = \mathbb{P}\left(X \leq \frac{y}{2}\right) = F_X\left(\frac{y}{2}\right)$$

We then obtain the pdf for Y by differentiation and the chain rule.

$$f_Y(y) = F'_Y(y) = \underbrace{\frac{1}{2} F'_X\left(\frac{y}{2}\right)}_{\text{chain rule}} = \frac{1}{2} f_X\left(\frac{y}{2}\right) = \begin{cases} \frac{1}{2} & \frac{y}{2} \in [0, 1] \\ 0 & \text{otherwise} \end{cases}$$

Proposition 3.11 - Increasing Functions

Let X be a random variable and $Y := g(X)$ where $g : \mathbb{R} \rightarrow \mathbb{R}$ is a strictly increasing function. Then, g is invertible on its range (denoted g^{-1}). Thus the cdfs for X and Y are related as

$$F_Y(y) = \underbrace{\mathbb{P}(g(X) \leq y)}_{\text{as } g \text{ is increasing}} = \mathbb{P}(X \leq g^{-1}(y)) = F_X(g^{-1}(y))$$

Differentiating this gives us the pdfs

$$f_Y(y) = f_X(g^{-1}(y)) \frac{dg^{-1}(y)}{dy} = \left. \frac{f_X(x)}{g'(x)} \right|_{x=g^{-1}(y)}$$

The probability mass of X in the interval $(x, x + dx)$ gets mapped to the interval $(g(x), g(x + dx))$. By *Taylor Expansion* of g we get $g(x + dx) \simeq g(x) + g'(x)dx$

Proposition 3.12 - General Mappings - Single Random Variable

Let X be a (discrete) random variable and define $Y := g(x)$ for any $g : \mathbb{R} \rightarrow \mathbb{R}$ which is differentiable.

There is a contribution $\frac{f_X(x)}{|g'(x)|}$ from each x such that $g(x) = y$. Where $g'(x)$ is positive or negative does not matter, as it only determines where the pre-image of $(y, y + dy)$ is of the form $(x, x + dx)$ or $(x - dx, x)$. Only the relative widths of the intervals matters for the contribution.

By summing the contributions from all solutions of $g(x) = y$ we get

$$p_Y(y) = \sum_{x:g(x)=y} \frac{f_X(x)}{|g'(x)|}$$

This formula is valid so long as the set $\{x : g(x) = y\}$ is countable. If it is not countable, then Y is a continuous RV (or is a mixed random variable)

Proposition 3.13 - General Mappings - Random Vectors

Consider the random vector $\mathbf{X} := (X_1, \dots, X_d)$ with joint density $f_{\mathbf{X}}$ and define $\mathbf{Y} := g(\mathbf{X})$ for any differentiable $g : \mathbb{R}^d \rightarrow \mathbb{R}^d$.

Let $J_g(\mathbf{x})$ denote the *Jacobian* of g at \mathbf{x} . Then

$$f_{\mathbf{Y}}(\mathbf{y}) = \sum_{\mathbf{x} \in \mathbb{R}^d : g(\mathbf{x}) = \mathbf{y}} \frac{f_{\mathbf{X}}(\mathbf{x})}{|\det(J_g(\mathbf{x}))|}$$

Let \mathbf{x} solve $g(\mathbf{x}) = \mathbf{y}$. In a neighbourhood of \mathbf{x} , g is approximately a linear function. By *Taylor Expansion*

$$g(\mathbf{x}') \simeq f(\mathbf{x}) + J_g(\mathbf{x})(\mathbf{x}' - \mathbf{x}) = \mathbf{y} + J_g(\mathbf{x})(\mathbf{x}' - \mathbf{x})$$

for \mathbf{x}' in a small enough neighbourhood of \mathbf{x} .

0 Reference

Definition 0.1 - Stochastic Matrix

A matrix is called a *Stochastic matrix* if:

- i). All elements are non-negative.
- ii). All rows sum to 1.


Definition 0.2 - Convex Function

A function $f : \mathbb{R} \rightarrow (\mathbb{R} \cup \{+\infty\})$ is *Convex* if, $\forall x, y \in \mathbb{R}, \alpha \in [0, 1]$, we have

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

A smooth function f is convex iff f is twice differentiable and $f''(x) \geq 0 \forall x \in \mathbb{R}$.

Visually, a function is convex if you can draw a line between any two points on the function and the function lies below the line.



Definition 0.3 - Equivalence Relation

A relation is an *Equivalence Relation* if it is

- i). Reflexive: $i \leftrightarrow i$.
- ii). Symmetric: If $i \rightarrow j$ then $j \rightarrow i$.
- iii). Transitive: If $i \rightarrow j$ and $j \rightarrow k$ then $i \rightarrow k$.

Definition 0.4 - Simple Random Walk

A *Simple Random Walk* is a random walk which moves only one step at a time. (i.e. $X_{t+1} = X_t \pm 1$). A probability p is defined for $\mathbb{P}(X_{t+1} = X_t + 1)$, this means $1 - p$ is the probability of stepping in the other direction. A *Simple Random Walk* is *Assymmetric* if $p \neq 1 - p$.

Definition 0.5 - Matrix Exponential, e^X

Let X be a matrix then we define the *Matrix Exponential* as $e^X := I + X + \frac{X^2}{2!} + \dots$

Theorem 0.1 - Pinsker's Inequality

For two distributions $\text{Bern}(p)$ and $\text{Bern}(q)$

$$K(q; p) \geq 2(q - p)^2$$

Definition 0.6 - Jacobian Matrix

The *Jacobian Matrix* is the first-order partial-derivatives of a multidimensional function $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ wrt each parameter.

$$J_f := \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix}$$

0.1 Notation*Multi-Armed Bandit Problem*

$p_{x,y}$	$\mathbb{P}(X_{t+1} = x X_t = y)$ for a <i>time homogenous markov process</i> .
$I(t)$	Arm played in round t .
$N_i(t)$	$\sum_{s=1}^t \{I(s) = i\}$ Number of times arm i was played in first t rounds.
$S_i(t)$	$\sum_{s=1}^t X_i(s) \{I(s) = i\}$ Total reward from arm i in first t rounds.
$\hat{\mu}_{i,N_i(t)}$	$\frac{S_i(t)}{N_i(t)}$ sample mean reward from arm i in first t rounds.
Δ_i	$\mu^* - \mu_i$ the arm gaps from a K -armed bandit.
$X_i(t)$	RV modelling the result if arm i was played on round t .

Markov Decision Problems

T	Time-Horizon.
S	State-Space.
A	Action-Space.
$A(s)$	Available actions when in system state s .
Y_t	Random variable modelling action taken by the agent at epoch t .
$r_t(s, a)$	The value of the reward received at epoch t if the system is in state s and the agent takes action a .
$p_t(s' s, a)$	The probability that the system is in state s' at time $t + 1$ given at time t the system was in state s and the agent takes action a .
$d_t(\cdot), q_t(\cdot)$	The decision function at epoch t .
π	A decision policy.
$HR(T)$	The set of policies on time-horizon T in which the decision rules are HRs.
$MD(T)$	The set of policies on time-horizon T in which the decision rules are MDs.
$\mathbb{P}^\pi, \mathbb{E}^\pi$	Probability and expectation given policy π is being used.
T^π	The expected reward for a given policy.
$u * (s)$	The terminal reward for a given MDP.

General

$z_{m:n}$	The m^{th} to the n^{th} elements of a sequence $\{z_t\}$. height
-----------	--

0.1.1 Asymptotic Notation**Definition 0.7 - Oh Notation**

Let $f, g : \mathbb{R}^+ \rightarrow \mathbb{R}$.

We say $f = o(g)$ (little oh of g) at 0 if $\frac{f(x)}{g(x)} \xrightarrow{x \rightarrow 0} 0$.

We say $f = O(g)$ (big oh of g) at 0 if $\exists c > 0$ st $|f(x)| \leq c|g(x)|$ in a neighbourhood of 0.

$f = o(g)$ at infinity and $f = O(g)$ at infinity are defined analogously.

Definition 0.8 - Omega Notation

Let $f, g : \mathbb{R}^+ \rightarrow \mathbb{R}$.

We say $g = o(f)$ if $g = o(f)$ and we say $f = \Omega(g)$ if $g = O(f)$.

Example 0.1 - Oh & Omega Notation

Define $f(x) = x$, $g(x) = \sin(x)$ and $h(x) = x^2$.

Then, $g = O(f)$ at 0 and $g = o(f)$ at infinity. $h = o(f)$ at 0 and $h = \omega(f)$ at infinity.